

When Protocol States Mislead: Evidence-Gated Admission and Online Feedback Calibration for LLM-Assisted Stateful Fuzzing

Ketang Chen^a, Xiaolei Ren^{a,*}, Junxi Chen^a

^aMacau University of Science and Technology, Avenida Wai Long, Taipa, Macau, China

Abstract

Response-derived states offer inexpensive feedback for stateful protocol fuzzing, but state novelty is an imperfect proxy for code progress. Requests generated by large language models (LLMs) can amplify this mismatch when unproductive proposals receive durable mutation resources. We present LoopFuzz as an integrated controller design and a retrospective boundary analysis. Bounded trials distinguish structural validity, response acceptability, and reachability from code and inferred protocol state machine (IPSM) novelty. The design assigns code-supported candidates to durable admission and state-only candidates to provisional validation. Discounted Beta–Bernoulli updates and Thompson sampling provide an exploratory mechanism for online state-productivity calibration. Analysis of 642 experimental records across nine targets yields descriptive coverage, cost, admission, and proxy-reward measurements. Observed code-branch and IPSM state-edge endpoints do not consistently reproduce the hypothesized mismatch across motivating targets. Four vulnerability cases provide replay evidence with differing patch support. Incomplete candidate lifecycles, heterogeneous observation horizons, and uncertain operating conditions prevent causal attribution to admission or calibration and do not establish controlled vulnerability rediscovery. The study distinguishes the proposed evidential criteria from the measurements needed to validate the design.

Keywords: Stateful protocol fuzzing, Large language models, Response-derived state feedback, Evidence-gated admission, Online productivity calibration

1. Introduction

Stateful protocol fuzzing must discover useful interaction sequences as well as useful message contents. A server can accept a request only after earlier exchanges establish a session, authenticate a client, or create a resource. AFLNet uses responses to construct an inferred protocol state machine (IPSM), making protocol-visible feedback available without a target-specific internal-state extractor [1]. Limited target-specific knowledge introduces an abstraction boundary: a response code or sequence is evidence about observable behavior, rather than an identifier of the server’s complete semantic state.

StateAFL already identifies limitations of response-derived state abstractions and uses memory observations to infer server states [2]. The fact that responses can misrepresent internal progress is therefore established background. Our question is how a fuzzer retaining a response-visible abstraction should decide, during a campaign, whether a particular state signal is useful enough to justify further exploration. A new IPSM state edge can accompany a previously unseen code branch, reproduce an already explored rejection path, or distinguish responses after measured code coverage has saturated. These outcomes should not receive identical evidential weight.

Assistance from large language models (LLMs) makes this distinction consequential for queue management. ChatAFL demonstrates that LLMs can derive protocol structures and propose interaction sequences [3]. A generated request is nevertheless a hypothesis about both message structure and session state. If a proposal is retained indefinitely because it produces a new response-derived state, the fuzzer may repeatedly mutate an input whose apparent novelty has little downstream code value. Durable queue membership represents a continuing allocation of execution resources.

Three target families motivate the investigation. Forked-daapd exposes DAAP/DACP behavior through HTTP status responses, so response distinctions may overstate useful code progress. Lighttpd1 can execute different code paths under the same HTTP status, so response novelty may understate code progress. LightFTP motivates examining continuing IPSM growth after code-coverage saturation. These hypotheses motivate the evaluation; they are not assumed to hold in the present observations. Section 2 separates them from the available measurements.

LoopFuzz is studied here as an integrated controller design with a retrospective boundary analysis. Its intended admission sequence checks proposal structure, executes a bounded trial, and separates response acceptability, reachability, code novelty, and IPSM novelty. The design assigns code-supported candidates to durable admission and state-only candidates to provisional validation. An exploratory scheduling component uses

*Corresponding author: Xiaolei Ren.

Email addresses: 3250006913@student.must.edu.mo (Ketang Chen), x1ren@must.edu.mo (Xiaolei Ren), 2240016952@student.must.edu.cn (Junxi Chen)

discounted productivity estimates to relate response-state selection to subsequent code progress. The observations describe these signals without establishing full adherence to the decision criteria in every run.

The study contributes a controller specification and descriptive evidence about its validation boundaries:

1. **An integrated evidence-gated controller design.** We specify execute-before-promote admission with reject, provisional, and durable outcomes. The design separates IPSM novelty from code evidence and bounds provisional validation. Productivity calibration is an exploratory scheduling component; Beta–Bernoulli estimation and Thompson sampling are reused methods.
2. **A boundary analysis of response-derived feedback.** We distinguish independent code progress from IPSM novelty and examine when recorded admission and reward signals fall short of the intended decision criteria.
3. **A retrospective empirical study.** We report descriptive coverage, cost, admission, and proxy-reward diagnostics across nine targets, together with four vulnerability replay cases. The analysis preserves configuration uncertainty and separates replay evidence from controlled rediscovery.

These contributions concern the controller design and its observed evidence boundaries. They do not establish reduced queue pollution, improved downstream code productivity, or improved vulnerability rediscovery. The five-arm framework specifies the controls needed to test such claims.

2. Background and Motivation

2.1. Response-visible state and independent code feedback

Let $\mathcal{M}_t = (Q_t, \widehat{\Delta}_t, q_0)$ be the IPSM observed by time t , where Q_t contains response-derived state identifiers, $\widehat{\Delta}_t \subseteq Q_t \times Q_t$ contains observed IPSM state edges, and q_0 denotes the initial connection state. The abstraction can alias different internal states and separate executions that reach the same code. It can also depend on history, asynchronous replies, and reset behavior. It is not assumed deterministic or semantically complete.

Let $\mathcal{B}_t$ be the covered code-branch set under a fixed coverage instrumentation and replay procedure. An IPSM state-edge increment changes $\widehat{\Delta}_t$; a code-branch increment changes $\mathcal{B}_t$. These objects have distinct units. An AFL bitmap observation must also distinguish a newly encountered code edge from a hit-count bucket change. Neither bitmap slots nor response-state identifiers are interchangeable with source-level code branches.

2.2. Three falsifiable mismatch cases

Table 1 defines the three motivating cases. Figure 1 summarizes the observed A/D endpoints as mean $\pm$ sample standard deviation (SD), with actual run counts n and nominal $N = 10$. Establishing a mismatch mechanism requires comparable executable versions, a shared observation horizon, and execution-level or time-resolved evidence. Endpoint differences alone

do not establish which state caused a difference. Contradictory observations are retained when assessing the motivating hypotheses.

The desired controller must accommodate both overestimation and underestimation. Rejecting every state-only candidate could lose useful bridges; promoting every such candidate would treat an imperfect proxy as sufficient evidence for permanent resource allocation. A provisional queue expresses the intermediate conclusion that an input deserves limited descendant validation while its longer-term value remains unproven.

The scope covers textual or semi-textual request-response protocols with observable structure, responses, and code coverage. Response-derived inference requires no internal-state access, but code-based productivity makes the controller grey-box rather than purely black-box. Encrypted or opaque binary protocols require additional assumptions about request construction, response interpretation, and code observability.

3. Related Work and Novelty Boundary

Stateful greybox fuzzing studies how to expose, infer, and schedule protocol state [1, 2, 4, 5]. AFLNet supplies a response-derived IPSM; StateAFL uses memory snapshots, and NSFuzz traces program state variables. The subsequent AFLNet study revisits coverage-guided protocol fuzzing [6], while greybox protocol-state learning investigates how internal observations inform inferred models [7]. LoopFuzz retains the response-visible interface and studies its instrumental reliability under an independent code-progress criterion. It does not claim to originate state inference or to first identify misleading responses.

Execution control and feedback interpretation address complementary problems. Nyx-Net uses incremental snapshots for network fuzzing [8]; SnapFuzz targets high-throughput execution of network applications [9]; fault-injection-based network fuzzing provides another approach to exercising protocol interactions [10]. BLuEMan simulates interactions between actual Bluetooth Low Energy (BLE) stacks [11]. DistFuzz combines events, faults, and timing with message-sequence feedback and symmetry-based pruning for distributed systems [12]. These studies motivate treating execution conditions and observable feedback as separate design choices. Their protocol domains and execution mechanisms are not evidence that LoopFuzz generalizes beyond its evaluated servers.

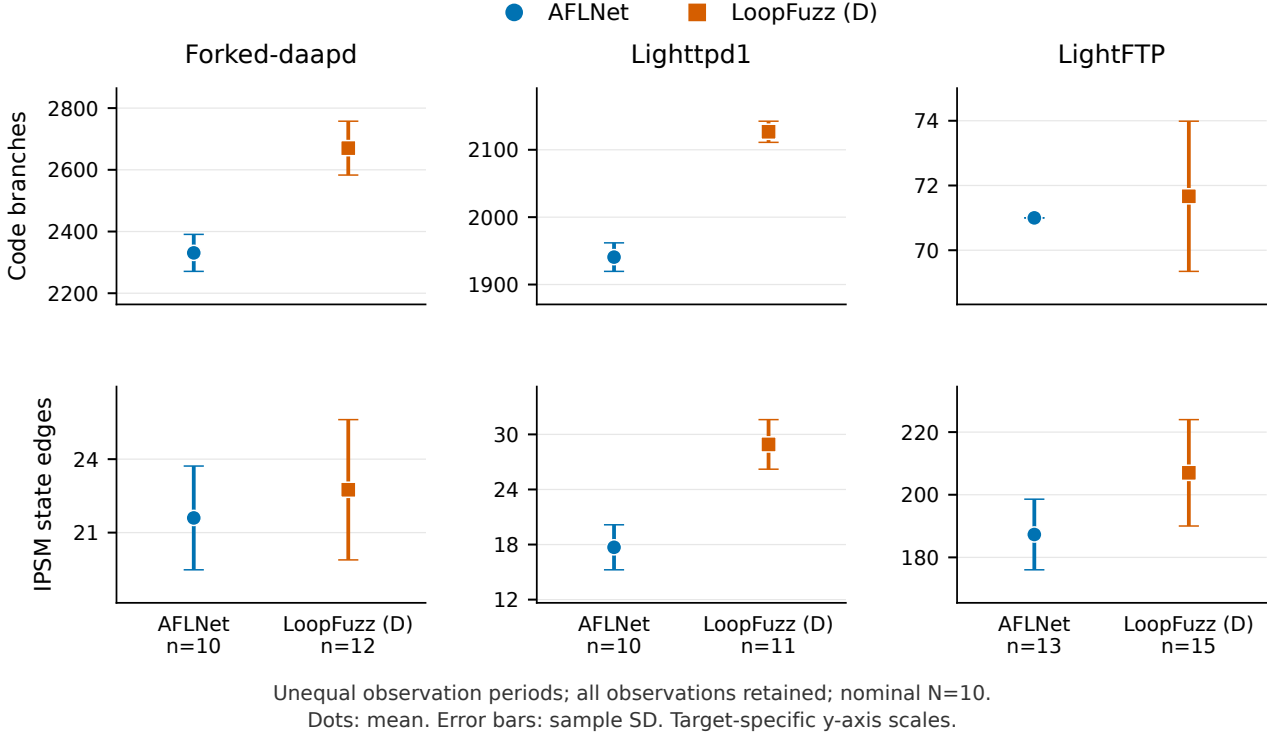
Structured generation reduces the difficulty of producing admissible inputs. Gramatron generates inputs through grammar automata [13]; CarpetFuzz extracts option constraints from documentation [14]. FormatFuzzer derives parsers, mutators, and generators from binary-format specifications [15]; Fuzztruction instead injects faults into a generating program to retain implicit format knowledge [16]. MINER uses valid sequence templates and learned request parameters to guide later requests [17]. These approaches concern how candidates are constructed. Our admission rule asks a subsequent question: whether a candidate has enough independently observed code value to acquire durable queue resources.

ChatAFL uses LLMs for protocol-message generation [3]; LLMIF derives device-testing inputs from protocol specifica-

Table 1: Motivating hypotheses and current evidence.

Target	Hypothesis	Required observation	Current interpretation
Forked-daapd	Overestimation: IPSM-only growth.	IPSM growth; low code reward.	Code regression not reproduced.
Lighttpd1	Underestimation: weak state signal.	Code gain with little IPSM growth.	Both endpoints increase; mechanism unverified.
LightFTP	Saturation: IPSM growth after code plateau.	Aligned code/IPSM saturation trajectories.	Near-equal code endpoints; matched-version trajectories needed.

Observed code and IPSM endpoints

Figure 1: A/D code-branch and IPSM state-edge endpoints (mean $\pm$ SD; actual n).

tions [18]; Fuzz4All generates and mutates program inputs across languages [19]. FuzzGPT studies unusual programs for deep-learning libraries [20], and WhiteFox uses compiler source information to guide optimization-triggering tests [21]. ProphetFuzz predicts high-risk option combinations from documentation [22]. Recent studies also address LLM-assisted hybrid fuzzing [23] and LLM-assisted directed greybox fuzzing [24]. These varied generation objectives reinforce the distinction between a model’s proposal and evidence obtained by executing it. The comparison specified here preserves ChatAFL’s scheduling lineage. Admission and calibration require separate controlled contrasts.

Adaptive resource allocation is established in fuzzing. The Bandit’s States models protocol-state selection as a bandit problem [25]. T-Scheduler combines Beta–Bernoulli estimation with Thompson sampling for seed scheduling [26]. Entropic allocates seed energy using estimated information gain [27]. SeqFuzzSDN infers extended finite-state machines and plans control-message sequences for software-defined networking con-

trollers [28]. These works delimit the contribution. FOX formulates coverage-guided fuzzing as online stochastic control [29], while FISHFUZZ dynamically prioritizes targets and seeds using distance information [30]. WingMuzz schedules open-source protocol programs and their seeds to guide blackbox testing [31]. Thus neither adaptive selection nor a productivity score is novel by itself. The distinction here is the combination of a response-derived proxy, an independent code-progress target, provisional and durable LLM-candidate admission, and traceable episode-level evidence.

Protocol correctness also requires an oracle beyond apparent acceptance. ProtocolGuard combines specification-derived checks with dynamic verification of non-compliance [32]; BSFuzzer uses context-aware semantic testing for BLE logic flaws [33]; MerCuriuzz targets logical vulnerabilities in QUIC systems [34]. These studies examine semantic violations, whereas the proposed calibration reward targets code productivity. A response-state transition or positive code reward therefore does not certify security impact. Magma’s ground-truth bug

Table 2: Prior mechanisms and novelty boundaries.

Prior work	Established focus	Present boundary
AFLNet; StateAFL; NSFuzz ChatAFL; LLM synthesis	Response-derived or internal-state feedback. Protocol-guided message/sequence generation.	Proxy calibration, not semantic-state recovery. Execute-before-promote; bounded provisional validation.
The Bandit’s States [25]	Bandit-based protocol-state selection.	Couple proxy calibration with admission.
T-Scheduler [26]; Entropic [27]	Posterior- or entropy-based seed scheduling.	Calibrate response proxies.
SeqFuzzSDN [28]	Learned state models; message-sequence planning.	Response-proxy productivity; candidate admission.
Magma [35]; evaluation studies	Ground-truth bug oracles; repeated trials.	Information-controlled network rediscovery; minimal security-patch pairs.

conditions motivate distinguishing bug-related reachability from triggering [35]; our historical vulnerability protocol additionally requires patch validation and controlled information exposure.

Evaluation methodology is a further boundary. ProFuzzBench and FuzzBench provide structured evaluation settings for protocol and general-purpose fuzzers [36, 37]. Reliability analyses and prudent evaluation practices require care when interpreting coverage, repeated trials, and resource controls [38, 39]. Our protocol adopts these concerns while keeping available retrospective observations distinct from the unfinished confirmatory comparisons. Candidate-level, episode-level, and controlled rediscovery measurements remain necessary to establish the proposed mechanisms.

4. Evidence-Calibrated Controller

This section defines the admission and calibration criteria. Section 5 examines whether the available evidence establishes that these criteria were satisfied. A decision rule alone does not establish experimental adherence.

4.1. LLM-generated proposals

The model receives permitted RFC fragments, seed traces, and observed interaction summaries and returns candidate requests or bounded scheduling hints. Each candidate is linked to its source state, intended target/frontier, prompt, and model response. Message-field hypotheses remain proposals rather than verified protocol grammars.

The information policy specifies permitted documentation and observed responses, excluding CVE identifiers, security advisories, exploit inputs, and patch contents. Adaptive campaigns may produce different prompts despite identical information policies; prompts and responses must therefore remain traceable.

Request proposals undergo admission; scheduling proposals require structural, state-consistency, and seed-consistency checks and cannot insert durable inputs. Parsing or transport failure does not create a positive admission decision. Call/token totals include all attempts, retries, failures, and successful replies.

4.2. Admission predicates and evidence separation

For candidate x , $P(x)$ checks the prescribed format, message structure, required fields, lengths, and terminators before execution. A bounded trial then evaluates response acceptability

$U(x)$ and reachability $R(x)$: arrival at a declared target/frontier or escape from a recorded rejection sink. Table 3 gives observable evidence and decision criteria.

HTTP/FTP 4xx/5xx classes are rejection priors, not a universal semantic definition. Protocol-specific exceptions are pre-specified; responses and justifications accompany each decision. Timeout, reset, disconnect, or an unparseable response gives no positive U evidence. A valid response alone is insufficient for R . Evidence of security violations is retained independently even when normal request admission fails.

Gain is split into

$$G_{\text{code}}(x) = \mathbf{1}\{\text{new code coverage or a favored code-coverage entry}\}, \quad (1)$$

$$G_{\text{state}}(x) = \mathbf{1}\{\text{new IPSM node, state edge, or path}\}. \quad (2)$$

Favored status must be grounded in code coverage; merely being queued or state-novel does not qualify. Source-level code branches are the primary coverage measure. A code-edge reward requires a prespecified, separately reported definition. Hit-count novelty is a separate sensitivity condition. Pre-trial coverage and IPSM observations identify new gains.

4.3. Provisional and durable queues

Let Q_{prov} contain temporary candidates and Q_{dur} contain long-term mutation seeds. Figure 2 summarizes the three admission outcomes:

$$D(x) = \begin{cases} \text{durable,} & P \wedge U \wedge R \wedge G_{\text{code}}, \\ \text{provisional,} & P \wedge U \wedge R \wedge \neg G_{\text{code}} \wedge G_{\text{state}}, \\ \text{reject,} & \text{otherwise.} \end{cases} \quad (3)$$

IPSM novelty alone is never sufficient for durable promotion. These logical conditions do not assume statistical independence among predicates.

A provisional entry receives at most $H = 64$ descendant executions and a local validation deadline of 30 seconds, whichever is exhausted first; the proposed provisional capacity is 64. The design requires pilot assessment and parameter specification before formal runs. Budget checks occur before each descendant execution. Budget exhaustion removes the entry from active scheduling while preserving its supporting evidence. Campaign termination before the validation horizon is censoring, not unproductive expiration.

Table 3: Admission predicates and required evidence.

Predicate	Pass condition	Failure	Measurements
P	Valid required structure.	Missing fields; invalid lengths/delimiters.	Candidate; criteria; reason; hypothesis version.
U	Accepted response.	Rejection; timeout/reset/drop; unparseable reply.	Response; acceptance rule; failure/exception.
R	Target/frontier reached or sink escaped.	Sink retained; destination missed.	Source/target; observed state sequence.
G_{code}	New code coverage or code-supported favored entry.	IPSM-only novelty.	Pre/post coverage; novelty; favored reason.
G_{state}	New IPSM node, state edge, or path.	Known IPSM structure.	Pre/post IPSM; new state transitions.

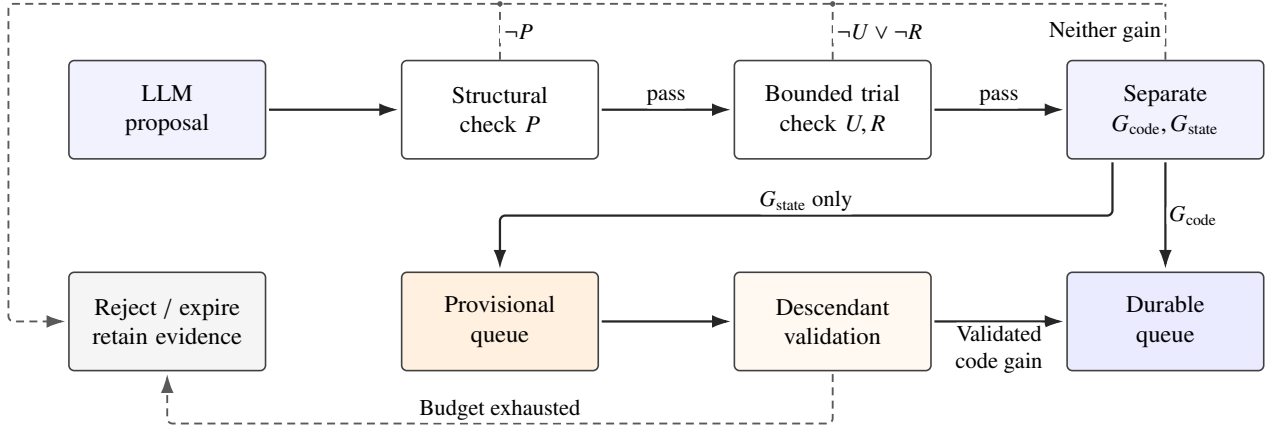


Figure 2: Evidence-gated admission and two-tier queues.

A descendant supplying qualifying code evidence permits promotion of the replayable parent sequence and productive descendant, subject to the same structural, response, and reachability checks. The observations identify the first productive descendant and the candidate promoted. Each code discovery is attributed once to its execution and candidate ancestry.

Productivity estimates select where to propose and continue exploration; they do not independently relax Equation 3. Fixed and calibrated arms use identical admission rules and provisional budgets. Calibration thus affects admission outcomes indirectly through exploration opportunities, preserving the causal scheduling comparison.

4.4. State-selection episodes and independent reward

An episode begins when the scheduler selects response-derived state s and a seed capable of reaching it. It receives one fixed execution-energy budget. Individual executions within the episode do not generate repeated posterior updates. Mutation counts, reset failures, elapsed time, and budget exhaustion are measured. Episodes interrupted by experimental environment failure are incomplete and receive no assigned reward.

For completed episode t targeting s_t , let ΔB_t count previously unseen code branches discovered within that episode. The main reward is

$$r_t = \mathbf{1}\{\Delta B_t > 0\}. \quad (4)$$

A declared code-edge instrumentation variant is admissible, but IPSM novelty is excluded. Favored descendants are a secondary sensitivity reward. Equal energy prevents a state appearing more

productive merely because it receives more executions. Trial execution, ordinary mutation, and provisional validation remain distinct in reward attribution, preventing duplicate updates.

θ_s denotes the probability of positive code reward conditional on selecting s under the episode policy and current campaign phase. It is instrumental utility, not the probability that s is a true semantic state. U continues to mean response acceptability; utility is never denoted $U(s)$.

4.5. Discounted productivity estimates

Each state starts with $\alpha_s = \beta_s = 1$ and

$$\theta_s \sim \text{Beta}(\alpha_s, \beta_s), \quad \widehat{\theta}_{s,t}^{\text{pre}} = \frac{\alpha_{s,t}}{\alpha_{s,t} + \beta_{s,t}}. \quad (5)$$

For a completed selected-state episode:

$$\alpha_s \leftarrow 1 + \gamma(\alpha_s - 1) + r_t, \quad (6)$$

$$\beta_s \leftarrow 1 + \gamma(\beta_s - 1) + (1 - r_t). \quad (7)$$

The proposed discount is $\gamma = 0.995$, with sensitivity values $\{0.99, 0.995, 1.0\}$. Parameters are prespecified after the pilot and before inspecting formal results. At $\gamma = 1$, updates reduce to cumulative Beta–Bernoulli counts. For $\gamma < 1$, excess prior counts are discounted to respond to changing productivity as coverage saturates.

The term productivity posterior denotes this discounted working model. With discounting it is not the exact posterior of a stationary Bernoulli process. Discounting occurs on completed episodes for the selected state; elapsed time alone does not

discount unselected states. Episode-indexed updating and limited initial evidence therefore constrain applicability.

Every calibration metric uses the pre-update mean recorded before reward. Post-update predictions would leak the outcome. Thompson samples remain separate from the Bernoulli probabilities used for reliability assessment.

4.6. Calibrated scheduling

The topology term favors underexplored observed states:

$$F(s) = \begin{cases} 4, & s \text{ is newly reachable or } \deg^+(s) \leq 1, \\ 2, & 1 < \deg^+(s) \leq 3, \\ 1, & \deg^+(s) > 3. \end{cases} \quad (8)$$

Only reachable states with eligible seeds participate. $\deg^+(s)$ counts outgoing IPSM state edges, not code branches. Let $B(s)$ denote the common base weight of C, D, and E. Define

$$\text{Frontier}(s) = B(s)F(s). \quad (9)$$

Keeping this factor and the initial exploration policy fixed prevents calibration from simultaneously changing the baseline scheduler. Adapting posterior-guided randomization [40] to the discounted, frontier-weighted controller, draw

$$\tilde{\theta}_s \sim \text{Beta}(\alpha_s, \beta_s) \quad (10)$$

and calculate

$$S(s) = \text{Frontier}(s)[\epsilon + (1 - \epsilon)\tilde{\theta}_s], \quad \epsilon = 0.1. \quad (11)$$

Select a state with probability proportional to $S(s)$, recording the pseudorandom seed. This is Thompson-sampled weighted selection, not an argmax policy. Under a finite eligible set, the positive floor gives nonzero selection probability, not a fixed number of selections.

The fixed-policy comparator retains the original penalty:

$$S_D(s) = \text{Frontier}(s)\rho_{\text{fixed}}(s). \quad (12)$$

Its queue-yield statistic $p_s = D_s/(N_s + 1)$ counts interesting queued paths D_s per state selections N_s . It is not a calibrated Bernoulli probability and can reflect state-only novelty or unequal execution energy. E replaces only ρ_{fixed} with the calibrated factor in Equation 11, preserving the shared base, admission, prompts, triggers, repair settings, and caps. Score discretization and initial cyclic state selection must be prespecified and verified before comparison.

4.7. Intervention policy and optional repair

Plateau interventions follow prespecified detection, timing, and retry rules. The baseline trigger combines IPSM state-edge growth with a sequence of unproductive selections. This dependence on the state proxy is held constant across the causal comparisons; scheduling calibration does not itself validate the intervention trigger.

The confirmatory C/D/E design requires repair to be disabled. Counterexample-guided repair is an optional procedure without quantitative evaluation in this study. The available observations do not verify repair status in every run, and no effectiveness or cost benefit is attributed to it.

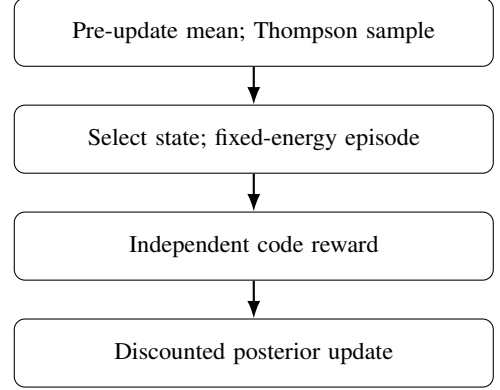


Figure 3: Online calibration from pre-update predictions and independent code rewards.

Algorithm 1: Evidence-gated candidate processing

Input: Candidate x , target, fixed pre-trial state

- 1 Associate candidate with proposal context and ancestry
- 2 **if** $\neg P(x)$ **then**
- 3 Reject x for invalid structure; stop
- 4 Execute one bounded trial without durable queue insertion
- 5 Assess $U, R, G_{\text{code}}, G_{\text{state}}$ and supporting evidence
- 6 **if** $P \wedge U \wedge R \wedge G_{\text{code}}$ **then**
- 7 Promote x to Q_{dur}
- 8 **else if** $P \wedge U \wedge R \wedge \neg G_{\text{code}} \wedge G_{\text{state}}$ **then**
- 9 Insert x into Q_{prov}
- 10 **while** budget, validation deadline, and campaign time remain **do**
- 11 Check budget before the next descendant execution
- 12 Execute one descendant and observe its outcome
- 13 **if** qualifying code gain and validation pass **then**
- 14 Identify first productive descendant; promote; stop
- 15 Classify provisional status: expired or censored
- 16 **else**
- 17 Reject x
- 18 Retain security-oracle evidence independently

5. Experimental Fidelity and Measurement

The evaluated system can represent code and state novelty separately, provisionally retain candidates, and condition state selection on productivity estimates. The available observations do not establish that these mechanisms operated as specified in every trial. Experimental fidelity requires verified conditions and admission order.

Candidate retention may precede behavioral verification. A subsequent U/R failure need not reverse that decision, so strict adherence to Equation 3 is not established. Coverage-based retention may also reflect execution-frequency changes rather than newly covered code branches. Similarly, evaluating provisional budgets only after an entire mutation cycle can exceed the prescribed limits.

IPSM novelty is assessed after changes associated with candidate retention. A candidate without prior code-based retention

Table 4: Decision criteria and evidence limitations.

Requirement	Current boundary	Evaluation consequence
Gate before promotion	Retention may precede U/R .	Queue protection unverified.
Independent state-only evidence	State assessment follows retention.	Zero provisional counts may be structural.
Complete P/R verification	P : recognition; R : transition.	Validity and reachability unverified.
New-code reward	Execution-frequency novelty may contribute.	Distinguish branches, frequency, and retention.
Fixed-energy episodes	Variable energy; zero executions.	Unequal opportunities.
Strict provisional budget	Checks after mutation cycles.	64-mutation/30-second limits may be exceeded.
Verified run conditions	Initial and effective settings may differ.	Within-run adherence requires verification.
Total call/token caps	64-call cap: plateau only; tokens unverified.	Total-budget fairness unverified.
Repair off in C/D/E	Repair status incompletely verified.	Arm labels do not prove adherence.

Table 5: Observation classes and required measurements.

Observation class	Experimental context	Measurements / outcomes
Experimental conditions	Arm; software/model versions; environment; sampling seed.	Call/token caps; resources; period; termination.
Candidate generation	Candidate/prompt; lineage; request/response; action; source/target.	Input/output tokens; generation time.
Admission trial	Candidate/trial; $P/U/R$; responses; reachability.	Code/IPSM changes; priority; latency; decision.
Provisional validation	Candidate; budget; first productive descendant.	Executions; code gain; conversion/expiration; latency; censoring.
State-selection episode	State/seed; pre-update posterior; sample; selection scores.	Mutations; code/IPSM gain; reward; updated posterior; duration/completion.
Security outcome	Anonymous oracle; reproducer; diagnostics; root cause.	Reached/triggered; first trigger; version-pair outcomes.
Repair (optional)	Lineage; counterexample; hypotheses; changed component.	Validation; post-repair predicates/gain; tokens; latency.

may therefore lack an independent state-only novelty assessment. This dependency can suppress the provisional condition by construction. An absence of provisional observations cannot distinguish effective filtering from a missing opportunity for provisional validation.

The available P decisions can reflect structural recognition without establishing full request validity. Likewise, R may indicate a state transition without confirming arrival at the intended target or frontier. These observations are weaker than the proposed predicates. Table 4 summarizes the resulting threats to experimental fidelity.

5.1. Measurement definitions and correspondence

Measurements distinguish experimental conditions, candidate generation, admission trials, provisional outcomes, scheduling episodes, and security outcomes; optional repair adds a seventh observation class. Table 5 specifies the corresponding evidence. Each observation identifies its trial, measurement definition, order, and time. CVE identities remain inaccessible during testing. Missing observations are not assigned zero values. Conditions distinguish target/fuzzer and model/service versions, sampling parameters, random seeds, and CPU/memory limits. Admission measurements retain reasons, state sequences, and separate pre/post code and IPSM counts. Episodes retain pre/post α, β , the pre-update mean, Thompson sample, and frontier/selection scores.

Confirmatory evaluation requires at least 99.9% completeness of candidate observations, an identifiable candidate for every trial, an identifiable episode for every posterior update,

and reconstructable promotion decisions. The denominator includes unsuccessful generation and interpretation attempts. Observation availability, correspondence across measurements, and agreement between reported decisions and actual queue membership are assessed separately.

6. Evaluation Scope and Validation Criteria

The present study is a retrospective, descriptive analysis of benchmark, ablation, and vulnerability-replay observations. The five-arm design and measurement definitions below state the conditions required for causal validation; they do not establish that a matched confirmatory evaluation was completed. Prior registration, prospective parameter selection, and completed pilot validation cannot be inferred from these records.

6.1. Research questions

1. **RQ1: State/code alignment.** How reliably does response-derived state novelty predict downstream code progress across targets and campaign phases? Outcomes include IPSM state-edge/code-branch alignment, episode productivity, calibration, and boundary cases.
2. **RQ2: Admission.** Does execute-before-promote reduce durable queue pollution and improve downstream productivity under matched proposal and scheduling conditions? The intended contrast is D–C.
3. **RQ3: Calibration.** Does online productivity calibration improve scheduling over a fixed response-state heuristic? The intended contrast is E–D.

Table 6: Required configurations for arms A–E.

Arm	LLM	Admission	Scheduler	Role
A: AFLNet	No	Original AFLNet	AFLNet	Non-LLM baseline.
B: Controlled ChatAFL	Yes	ChatAFL direct/open-loop	ChatAFL lineage	B–A: LLM-lineage comparison.
C: LoopFuzz-direct	Yes	Parseable to durable	Fixed LoopFuzz	Admission counterfactual.
D: LoopFuzz-gated-fixed	Yes	Evidence gate; provisional queue	Same fixed LoopFuzz	D–C: admission effect.
E: LoopFuzz-gated-calibrated	Yes	Identical to D	Productivity calibration	E–D: calibration effect.

4. **RQ4: Historical vulnerability rediscovery.** Does the controller improve reproducible rediscovery under independent reached/triggered oracles and patch validation?
5. **RQ5: Cost and applicability.** What execution-time, throughput, and LLM costs arise, and under which target characteristics does the method fail to help?

The directional hypotheses are improved downstream productivity for D over C and greater area under the code-branch coverage curve (AUC) for E over D. Regressions, null results, and uncertainty are reported per target rather than hidden in pooled means or replaced by increased IPSM state-edge counts.

6.2. Five arms and causal comparisons

C/D share model configuration, prompt construction, trigger, repair setting, scheduler, and budget policy; the intended intervention is admission. E/D share admission and differ only in state selection. Controlled ChatAFL retains its scheduling lineage; replacing it with the LoopFuzz scheduler would change the baseline. B–A is not an isolated admission comparison.

Proposal control has two scopes. Paired diagnostics evaluate the same candidate under direct and gated admission from an identical initial target and campaign state, with equal descendant budgets. This comparison supports a candidate-level counterfactual. Independent adaptive campaigns diverge, so equal policies do not guarantee identical prompts or model outputs. Replay can assess structural validation and admission decisions under matched initial conditions, but cannot establish equivalence to an independently evolving campaign.

6.3. Targets and experimental controls

The primary design contains nine servers: four FTP servers, Exim (SMTP), Live555 (RTSP), Kamailio (SIP), Forked-daapd (DAAP/DACP over HTTP), and Lighttpd1 (HTTP). Table 7 identifies the target protocols and experimental conditions. The benchmark specifies source versions, local communication conditions, and initialization procedures. The correspondence between these versions and every tested program remains unverified. In the table, yes/no denotes dictionary use. Reset evidence distinguishes a declared procedure from an unrecorded one; neither label verifies successful state restoration.

Runs use isolated execution environments, local-network communication, and AFL-based coverage measurement. Table 7 reports the declared source versions, seed protocols, dictionaries, communication settings, and reset evidence. These declarations do not independently establish the identity of every tested program or the effective reset conditions throughout each run.

Each causal comparison requires identical software versions, compilation settings, coverage instrumentation, seeds, replay procedures, and reset policies. The confirmatory design randomizes trial order within machine blocks and balances trials across machines and periods. Random seeds, resource allocations, measured CPU consumption, resource contention, and execution rates characterize each trial. Allocated resources do not establish equal consumption; initialization, model waiting, validation, mutation, and termination all contribute to experimental cost. Resource-efficient benchmarking also makes the cost of obtaining reliable comparisons an explicit concern [41]. The proposed analysis therefore distinguishes allocated CPU-hours, measured consumption, and actual search exposure.

The design allocates 10 independent 24-hour runs per target to each arm, plus 10 for C/D/E, reaching 20 each. For nine targets and one allocated CPU per run, the stages require 10,800 and 6,480 CPU-hours, totaling 17,280 before overhead/failures. These are planned resources, not measured usage. A four-target pilot covers Forked-daapd, Lighttpd1, LightFTP, and ProFTPD or Pure-FTPd with five arms, three runs, and four to six hours each. Pilot observations are excluded from formal inference.

An intention-to-run analysis accounts for every planned and initiated trial, including completion, timeout, target or fuzzer failure, experimental environment or measurement failure, exclusion, and repetition, with reasons linking repeated attempts. Low coverage never justifies exclusion. Only prespecified experimental environment failures permit repetition. Interrupted observations are not extended to 24 hours; they contribute failure or censoring outcomes according to the estimand, with sensitivity analysis. Restriction to completed trials does not constitute intention-to-run analysis.

6.4. LLM fairness and configuration

The confirmatory design requires equal policies and resource ceilings across LLM arms; actual calls and token usage may differ with adaptive trajectories. In Table 8, configuration values are supported by final observations for C/D/E and the sensitivity groups, while equivalent B observations are unavailable. Timeout, retry, seed, and prompt-template settings describe the prescribed procedure; their actual use is not independently verified for every run. Unresolved settings remain identified explicitly.

An output-token limit per request is distinct from a total campaign limit. A recorded token-limit value of zero does not establish equal-token control without evidence of the effective policy. Monetary costs require verifiable usage and applicable prices; unavailable costs remain unspecified.

The present cost analysis reports observed calls and tokens descriptively. Estimating code gain per candidate, per 1,000 total

Table 7: Target protocols and recorded test conditions.

Target	Protocol	Source version	Local endpoint	Seed protocol	Dictionary	Reset evidence
LightFTP	FTP	139af7c	tcp/2200	FTP	yes	Declared
bftpd	FTP	6.1 release	tcp/21	FTP	yes	Declared
ProFTPD	FTP	61e621e743346	tcp/21	FTP	yes	Declared
Pure-FTPd	FTP	10122d9f	tcp/21	FTP	yes	Declared
Exim	SMTP	d6a5a05b84	tcp/25	SMTP	yes	Declared
Live555	RTSP	live.2023.05.10	tcp/8554	RTSP	yes	Unrecorded
Kamailio	SIP	a2209018fb03d	udp/5060	SIP	no	Declared
Forked-daapd	DAAP/DACP	2ca10d9	tcp/3689	DAAP	no	Unrecorded
Lighttpd1	HTTP	9f38b63cae3e2	tcp/8080	HTTP	yes	Unrecorded

Table 8: LLM settings for C/D/E and sensitivity groups; B unverified.

Parameter	Recorded value	Comparison requirement
Model / version / service	codex-auto-review; cctq.ai	Same resolved version; name insufficient.
Temperature by call class	grammar 0.5; plateau 1.2	Equal temperature within each call class.
Top- p / presence/frequency penalties	top- p 1.0; penalties unrecorded	Explicit values or defaults.
Output tokens per attempt	4,096	Same cap; include truncation/failures.
Attempted-call / total-token caps	call cap 64; token-limit setting 0	All request classes and retries included.
Timeout / retry policy	120 s request, 30 s connect, 60 s low-speed abort; 2–5 retries by call class, 2 s apart; nonrecoverable service failures end further attempts	Same policy; include waits/failures.
Model seed / support	none specified in model requests	Seed-support status; determinism unassured.
Prompt corpus / template	no independent corpus/template version; prompts identifiable; template tied to fuzzer version	Same information policy; traceable exchanges.
Plateau cadence / threshold / reset	state-edge threshold 512; cadence unrecorded	Same intervention/state-growth criteria.

tokens, or at matched cumulative budgets would additionally require linked gain and complete usage measurements. These efficiency effects are not estimated here. Future comparisons must account for unsuccessful requests, retries, validation, and any repair costs; zero-denominator ratios remain undefined.

6.5. Mechanism and calibration metrics

Available results comprise recorded candidate and trial counts, first-failure distributions, disposition labels, trial latency, and exploratory episode diagnostics. Denominators accompany reported rates. The following descendant-based criteria define what would be required to validate admission quality; they are not measured outcomes of this study.

For a prospective evaluation, downstream productivity uses a fixed descendant horizon rather than admission gain. For a declared cohort C of N candidates enrolled before their outcomes are known, let $Y_i(H) = 1$ if a later descendant produces new code coverage within H executions, and 0 if the full horizon is observed without gain. With complete observation, define

$$\text{Productivity@ } H = \frac{1}{N} \sum_{i \in C} Y_i(H). \quad (13)$$

The proposed initial horizon is $H = 64$; no such outcome is estimated here. The admission trial is excluded. A positive outcome can be established early; a negative requires the complete horizon. With n_+ known positives and $n_?$ censored unresolved

candidates, cohort bounds are $[n_+/N, (n_+ + n_?)/N]$. A valid paired diagnostic would reserve the same full validation horizon for both controllers and retain incomplete observations in the denominator. Durable and provisional cohorts are distinct; post-promotion productivity requires a fresh horizon.

An avoided-pollution claim requires a paired C–D comparison: direct-arm durable candidates without downstream gain must be linked to their retention under gating. D-arm rejections alone do not establish avoided pollution. Estimating false negatives would require a predeclared random sample of rejected candidates, common initial conditions, descendant lineage, and recorded validation costs. These measurements are unavailable in the present study.

For complete episodes, $p_t = \widehat{\theta}_{s_t, t}^{\text{pre}}$ predicts later reward r_t . We assess probabilistic accuracy with the binary Brier score (BS), also used in software defect prediction [42]:

$$\text{BS} = \frac{1}{n} \sum_{t=1}^n (p_t - r_t)^2. \quad (14)$$

We adapt the binned expected calibration error (ECE) of Guo et al. [43] to binary episode productivity: within ten fixed equal-width bins, we compare the mean predicted reward probability with the observed reward frequency and weight each absolute gap by its episode count. This evaluates event probabilities rather than predicted-class confidence. Bin counts accompany the estimate; empty bins remain unestimated. Reliability plots and

precision–recall analysis [44] accompany the BS. We report non-interpolated average precision (AP): precision at each distinct score threshold weighted by its recall increment. This quantity differs from trapezoidal precision–recall area; it is left undefined when there are no positive episodes. A prequential base-rate predictor provides a BS reference because rare rewards can favor uninformative near-zero predictions. IPSM state-edge/code-gain ratios with zero code gain are undefined or shown as separate zero-gain cases.

Episodes within a run are dependent, so inferential uncertainty requires run-level clustering and target-specific analysis. Posterior trajectories and state weights remain descriptive diagnostics. A state/code association would require linked measurements within target and phase; pooled counts across unlike programs cannot establish that relationship.

6.6. Coverage and statistical analysis

The confirmatory outcomes are code-branch AUC over 24 hours and final code-branch count. AUC uses each run’s cumulative time series with a common interpolation convention and instrumentation, in branch-hours. Endpoints recorded at unequal durations are not substituted. Time to 95% of a target-specific common reference endpoint uses a threshold prespecified before formal comparisons; unreached thresholds are censored.

The confirmatory analysis specifies target-specific means, medians, and dispersion. Bootstrap resampling [45] operates on whole runs; the percentile 95% confidence intervals (CIs) and resampling hierarchy are analysis choices specified here. Planned independent run-level contrasts use two-sided Mann–Whitney U tests [46] with ties accounted for, and Cliff’s δ [47], oriented as $\Pr(X_{\text{treatment}} > X_{\text{comparator}}) - \Pr(X_{\text{treatment}} < X_{\text{comparator}})$. The rank test is not interpreted as a test of median differences alone. The primary family comprises 36 coverage tests across nine targets: D–C and E–D contrasts for final branches and AUC. The design specifies the Benjamini–Hochberg procedure [48] at nominal $q = 0.05$. Its original guarantee assumes independent test statistics; shared comparator runs and correlated coverage outcomes do not establish that condition here. Adjusted values would not imply unconditional false-discovery-rate control. A confirmatory admission-mechanism family would contain nine D–C productivity contrasts; these are not estimated here. Other diagnostics and sensitivity tests are exploratory. Effect direction is treatment minus comparator.

Confirmatory cross-target summaries would use equal-weight target effects and a hierarchical bootstrap resampling targets, then runs within targets. Raw branch counts from different programs or individual executions are not pooled as independent samples. Target regressions and uncertainty accompany aggregates. A future historical time-to-trigger analysis would use Kaplan–Meier estimates [49] under non-informative right censoring, alongside CVE-specific deadline recall. Fuzzer-failure interruptions are distinguished because censoring may be informative.

7. Available Evidence and Results

The analysis covers benchmark and ablation outcomes. Results obtained under earlier configurations do not establish the effectiveness of the proposed controller. Tables report observed quantities; unmeasured effects remain limitations rather than zero-valued results.

7.1. Sample composition and observation periods

The dataset contains 94 AFLNet (A), 91 ChatAFL (B), 104 benchmark LoopFuzz (D), 90 direct-admission (C), and 91 calibrated (E) observations. The sensitivity conditions contain 81 observations at $\gamma = 0.99$ and 91 at $\gamma = 1.0$. Settings reported at trial completion identify fixed-policy labels for 17 of the 90 nominal C observations and 12 of the 91 nominal E observations; the $\gamma = 0.99$ and $\gamma = 1.0$ groups contain 15 and 12 such records. Within C, all-durable, no-rejection outcomes corroborate direct admission. Posterior observations occur in every LoopFuzz condition and therefore do not distinguish E or the sensitivity groups. Nominal arms also contain multiple program versions: D four (78/12/9/5 runs), C two (83/7), and E two (71/20); each sensitivity group uses one version. Original group assignments are retained, but correspondence to effective conditions remains uncertain. All observations are included without assuming independence. The nominal target is ten trials per condition; means and sample SDs use observed values only, with actual n and no imputation.

All 104 benchmark LoopFuzz D observations are reported as completed; Table 9 gives target-specific counts. The interruption column identifies a reported termination category; it does not establish memory exhaustion or vulnerability activation. The reported duration measures active fuzzing rather than total CPU consumption or the entire experimental period.

The benchmark LoopFuzz observations also constitute the fixed-policy admission condition D and are counted once. The direct-admission and calibrated-scheduling conditions correspond to C and E. Two sensitivity groups use $\gamma = 0.99$ and $\gamma = 1.0$. The configuration uncertainty in Table 4 and unequal sample sizes preclude a fully matched comparison. Recorded posterior estimates alone do not establish a calibration effect.

The complete target–arm results, including both gamma sensitivity groups, are reported in Appendix A (Table A.1).

Table 10 includes every available A and D observation. Code-branch and IPSM state-edge endpoints represent the terminal measurements from each trial. For all 642 experimental records, the final trajectory measurement agrees with the corresponding endpoint. These are descriptive observations at unequal durations, rather than matched 24-hour outcomes. Sample SDs describe variation among the available observations; no causal effect or significance claim is attached to the endpoint differences.

7.2. RQ1: State–code alignment

Forked-daapd has mean code-branch endpoints of 2330.8 ± 60.0 for AFLNet and 2670.2 ± 87.3 for LoopFuzz, while IPSM state-edge endpoints are 21.6 ± 2.1 and 22.8 ± 2.9 . This direction does not support the motivating hypothesis of code regression.

Table 9: Observed trial counts and reported D-arm status ($N = 10$ nominal runs).

Target	AFLNet	ChatAFL	LoopFuzz(D)	Completed(D)	Interrupted(D)	Duration (min)
LightFTP	13	11	15	15	0	1473–1522
bftpd	10	10	10	10	0	1499–1510
ProFTPD	10	10	12	12	0	1489–1518
Pure-FTPd	11	10	10	10	0	1460–1472
Exim	10	10	10	10	0	1477–1486
Live555	10	10	12	12	0	1489–1505
Kamailio	10	10	12	12	0	1472–1518
Forked-daapd	10	10	12	12	0	1474–1524
Lighttpd1	10	10	11	11	0	1469–1523

Table 10: A/D terminal coverage (mean $\pm$ SD); descriptive comparisons.

Target	Code branches		IPSM state edges	
	AFLNet	LoopFuzz(D)	AFLNet	LoopFuzz(D)
LightFTP	71.0 $\pm$ 0.0	71.7 $\pm$ 2.3	187.3 $\pm$ 11.3	207.0 $\pm$ 17.0
bftpd	486.0 $\pm$ 8.8	502.7 $\pm$ 7.5	163.1 $\pm$ 15.4	226.7 $\pm$ 14.0
ProFTPD	5053.8 $\pm$ 97.4	5428.1 $\pm$ 104.0	190.5 $\pm$ 27.9	305.3 $\pm$ 20.3
Pure-FTPd	1076.4 $\pm$ 89.6	1223.3 $\pm$ 29.7	236.5 $\pm$ 26.7	321.4 $\pm$ 11.1
Exim	3798.8 $\pm$ 43.3	3914.9 $\pm$ 20.9	88.8 $\pm$ 5.2	152.7 $\pm$ 9.3
Live555	2893.6 $\pm$ 14.5	3067.1 $\pm$ 29.5	96.1 $\pm$ 8.6	170.1 $\pm$ 11.1
Kamailio	10002.6 $\pm$ 144.2	10631.2 $\pm$ 472.9	109.4 $\pm$ 9.7	128.2 $\pm$ 26.6
Forked-daapd	2330.8 $\pm$ 60.0	2670.2 $\pm$ 87.3	21.6 $\pm$ 2.1	22.8 $\pm$ 2.9
Lighttpd1	1940.7 $\pm$ 21.2	2126.7 $\pm$ 15.7	17.7 $\pm$ 2.5	28.9 $\pm$ 2.7

Table 11: D-arm candidate and episode counts; dispositions are reported labels.

Target	Candidates	Trials	Reject	Provisional	Durable	Episodes
LightFTP	48	48	44	0	4	13604
bftpd	13	13	13	0	0	3249
ProFTPD	0	0	0	0	0	2395
Pure-FTPd	173	173	166	0	7	5378
Exim	24	24	24	0	0	2650
Live555	851	851	844	0	7	7687
Kamailio	253	253	249	0	4	9004
Forked-daapd	412	411	410	0	1	12056
Lighttpd1	453	453	453	0	0	21570

The difference cannot be attributed to calibration, which is not established as active.

Lighttpd1 increases in both endpoint metrics: code branches from 1940.7 ± 21.2 to 2126.7 ± 15.7 , and IPSM state edges from 17.7 ± 2.5 to 28.9 ± 2.7 . These aggregate observations do not demonstrate the hypothesized small state expansion with hidden code progress. LightFTP has near-equal code endpoints, 71.0 ± 0.0 versus 71.7 ± 2.3 , with larger LoopFuzz IPSM state-edge counts, 187.3 ± 11.3 versus 207.0 ± 17.0 . That pattern motivates a saturation analysis but does not establish the full time course or identify individual misleading states.

Mechanism claims require matched versions and observation periods, with linked state/execution measurements, for all three cases. Figure 4 examines target-specific agreement between pre-update probabilities and binary rewards for coverage-based retention. All 32,946 E observations satisfy probability, Beta-mean, binary-response, uniqueness, temporal-order, and within-episode discounted-update consistency checks. Exclud-

ing 6,562 episodes with no mutations leaves 26,384 episodes from 91 runs. Three discontinuities between successive posterior estimates are retained; missing history is not reconstructed. These measurements do not establish policy adherence in every run.

Table 12 reports the BS, ten-bin ECE, and non-interpolated AP; lower BS/ECE is better. Its prequential reference uses only past included same-run rewards with Beta(1,1) smoothing. Percentile 95% CIs use 2,000 bootstrap resamples of runs within each target; dependence between trials remains unresolved. Empty bins are unestimated; CIs require at least two runs per bin. In Figure 4, points pair the bin-mean pre-update prediction with reward frequency; the diagonal denotes perfect calibration. Lower strips give positive-mutation episode counts, with actual runs n , episodes m , and nominal $N = 10$. Rewards measure coverage-based retention, potentially including execution-frequency novelty; nonzero mutation counts do not establish complete fixed-energy episodes. These diagnostics do

Table 12: Proxy-reward calibration in E (95% CIs in brackets).

Target	n	Episodes	BS	Reference BS	ECE (10 bins)	AP
LightFTP	10	6,622	0.188 [0.178, 0.198]	0.183 [0.172, 0.195]	0.166 [0.161, 0.172]	0.452 [0.430, 0.471]
bftpd	10	2,643	0.249 [0.243, 0.257]	0.247 [0.243, 0.252]	0.105 [0.094, 0.119]	0.517 [0.493, 0.542]
ProFTPD	10	1,002	0.183 [0.179, 0.188]	0.144 [0.126, 0.157]	0.137 [0.106, 0.174]	0.810 [0.781, 0.844]
Pure-FTPD	10	2,848	0.227 [0.217, 0.234]	0.243 [0.233, 0.248]	0.053 [0.041, 0.071]	0.700 [0.675, 0.722]
Exim	10	1,057	0.230 [0.221, 0.239]	0.219 [0.211, 0.227]	0.126 [0.106, 0.149]	0.703 [0.675, 0.734]
Live555	10	2,259	0.224 [0.212, 0.233]	0.218 [0.207, 0.229]	0.082 [0.067, 0.103]	0.742 [0.712, 0.771]
Kamailio	10	6,374	0.163 [0.153, 0.171]	0.161 [0.149, 0.173]	0.042 [0.032, 0.055]	0.858 [0.847, 0.866]
Forked-daapd	11	1,382	0.254 [0.245, 0.262]	0.253 [0.251, 0.256]	0.089 [0.070, 0.119]	0.596 [0.560, 0.623]
Lighttpd1	10	2,197	0.231 [0.225, 0.237]	0.251 [0.250, 0.253]	0.085 [0.072, 0.096]	0.665 [0.639, 0.688]

Pre-update probability and observed proxy reward

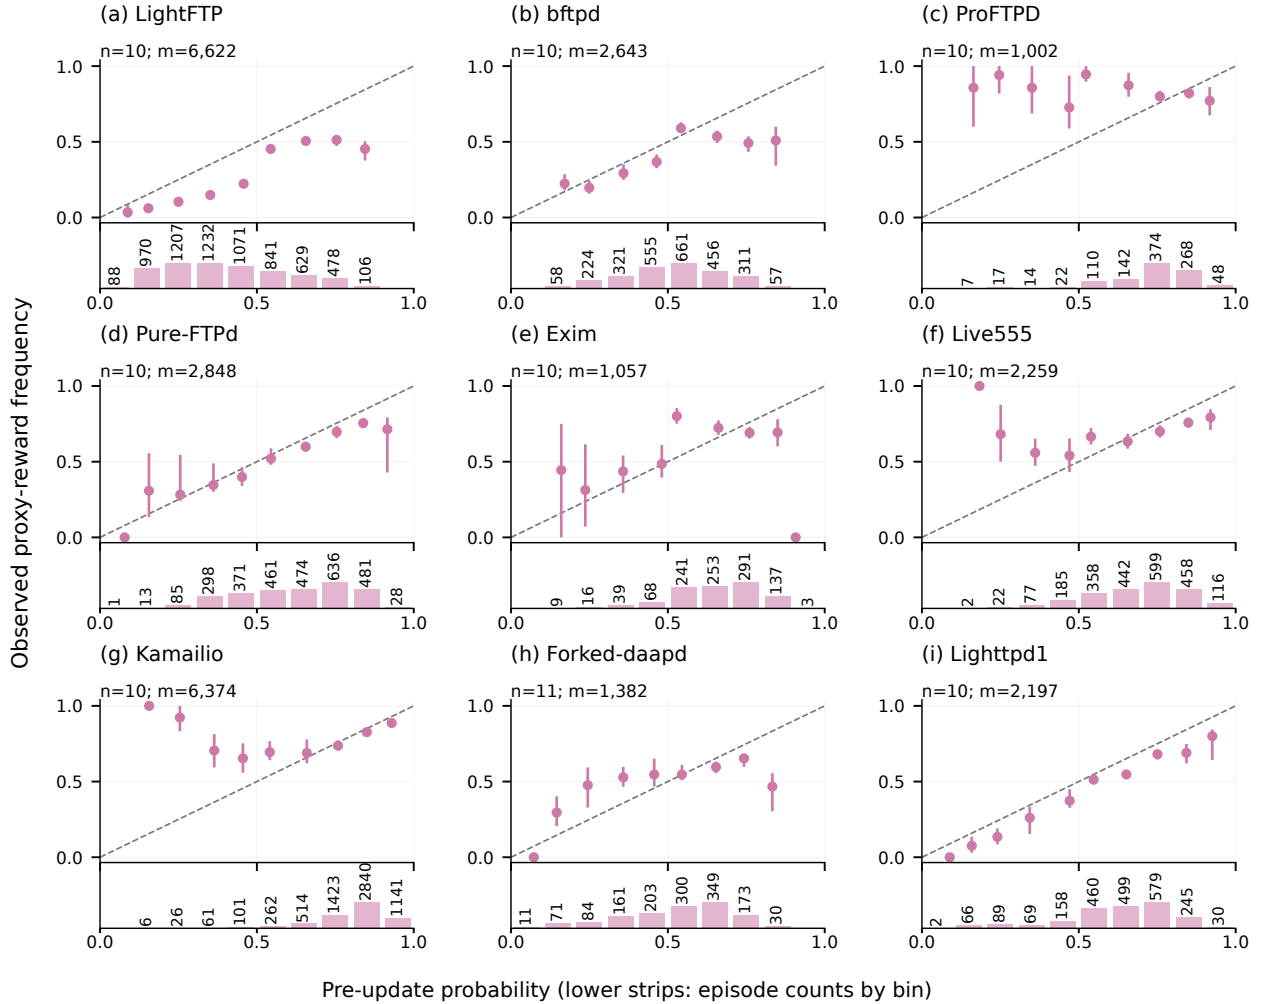


Figure 4: E-arm proxy-reward reliability (bin means; 95% CIs).

not establish independent code-branch calibration. The equal-energy subset in Table 14 conditions the same proxy reward on one energy and retains these limits.

7.3. RQ2: Admission records and evidence limits

The D condition contains 2,227 candidate observations and 2,226 validation trials, with 2,203 reported rejections, 23 durable admissions, and no provisional admissions. One candidate has

no corresponding trial; every trial has a corresponding candidate. C and E contain 1,326/1,326 and 746/745 candidates/trials, respectively. C has 1,326 reported durable admissions; E has 736 rejections and 9 durable admissions. These reported decisions do not independently verify queue membership or establish the required 99.9% observation completeness.

Candidate and admission observations are available for C, D, and E, with target-specific gaps. ProFTPD lacks these observations in D and E despite nonzero model usage. No provisional-lifecycle or repair-outcome observations are available across the 642 records. Reported provisional counts for admission, conversion, and expiration are zero throughout; the rates in Table 14 therefore describe reported totals, not independently verified queue behavior. Table 11 gives D-condition counts, and Table A.1 gives endpoint results. Missing observations establish neither zero conversion rates nor absence of provisional validation or repair.

Recorded dispositions do not independently verify durable queue membership. Without validated candidate-to-descendant links and a paired C–D comparison, promotion latency, downstream productivity, nonproductive retention, and admission false negatives cannot be estimated. Unreconciled usage also prevents attributing code gain to individual candidates or token expenditure. Accordingly, no reduction in queue pollution or improvement in downstream productivity or cost efficiency is claimed. Table 14 and Figure 7 report observed decisions and latency descriptively; first-failure classifications do not independently explain queue retention.

7.4. RQ3: Exploratory scheduling diagnostics

C, D, and E contain 56,164, 77,593, and 32,946 state-selection episodes, respectively. Mutation counts range from zero to 18,816 per episode, precluding an equal-energy interpretation of the full samples. Some coverage observations identify monitored execution locations or coverage-based retention decisions rather than independent source-level code branches. They therefore cannot be combined with terminal code-branch counts or interpreted as the primary reward in Equation 4.

An exploratory equal-energy subset holds the recorded mutation count constant for the corresponding diagnostics in Table 14; it does not establish matched target composition, run conditions, or a causal scheduling contrast. Restricting to episodes with exactly 265 mutations, the most frequent nonzero energy value shared by all three arms, leaves 3,986, 6,604, and 1,474 episodes for C, D, and E. Within this subset, mean recorded code edges per episode are 0.009, 0.003, and 0.016, and proxy-reward base rates are 0.004, 0.001, and 0.005: outcomes steeply rise with energy across all arms (episodes above 5,000 mutations are almost always rewarded), so this fixed low energy yields a subset with unequal sample sizes. Pre-update predictions are overconfident in this subset (BS 0.240, 0.192, 0.315; ten-bin ECE 0.428, 0.373, 0.516). The posteriors were updated under mixed energies; AP exceeds the recorded base rate (AP 0.012, 0.005, and 0.038 versus base rates near 0.005). These subset diagnostics are exploratory: they condition the observed proxy reward on one energy rather than measuring energy-conditioned calibration,

are pooled across targets without clustered uncertainty, and do not replace the per-target full-sample diagnostics in Table 12.

The evaluation includes the calibrated condition E and two discount-factor sensitivity conditions. Heterogeneous observation periods and uncertain effective scheduling policies prevent a matched causal interpretation of E–D. Table 13 reports the observed endpoints and AUCs; all 45 main target–arm combinations are represented, including ten E/bftpd observations. Settings reported at trial completion specify $\gamma = 0.99$ in 66 of the 81 $\gamma = 0.99$ observations and $\gamma = 1.0$ in 79 of the 91 $\gamma = 1.0$ observations; the remaining observations report a fixed policy at $\gamma = 0.995$. Four $\gamma = 0.99$ target groups remain below the nominal ten runs (Lighttpd1 6, Live555 7, Kamailio 9, Forked-daapd 9). A calibration benefit requires a prespecified analysis under matched observation periods.

Fig. 5 shows code-branch medians and interquartile ranges on a five-minute grid over 0–24 hours from the start of each run; the bands are not CIs. Linear interpolation stays within recorded observations, and labels give the contributing n at 24 hours. Table 13 instead integrates each full observed trajectory in branch-hours; cells show branches with actual n above AUC, both as mean $\pm$ SD. Each run contributes only between its first and last observation; time-dependent denominators accompany the figure data. Changes in the set of contributing runs can change an aggregate curve even when individual histories do not decrease. In 127 of the 470 observed trajectories, at least one recorded branch count decreases. The figure preserves these observations without enforcing monotonicity or otherwise altering the recorded measurements; both within-run decreases and changing support can produce aggregate decreases.

Fig. 6 presents three reproducibly selected state histories. For each motivating target, the E run is selected by the lower median terminal branch count, with ties resolved by a fixed ordering of observations; its two states with the most positive-mutation episodes are shown. Selection uses neither calibration accuracy nor effect direction. Complete histories, including zero-mutation episodes, show pre-update probabilities above the final-selection/frontier-score ratio. Time denotes episode end; lines guide the eye, and colors identify states within each target. Undefined zero-frontier ratios are omitted. These cases illustrate observed behavior without establishing a causal state/code mismatch.

7.5. RQ5: Recorded costs and interpretation limits

Table 15 reports model calls and token usage for D by target; supporting results also give C/E ranges and means. Model calls range from 52 to 248 per D run, so the stated limit of 64 does not establish a campaign-wide ceiling. Usage estimates have not been independently reconciled with service charges and may omit the final observation interval.

Settings reported at trial completion specify a model name, grammar/plateau temperatures of 0.5/1.2, top- p of 1.0, an output limit of 4096 tokens, a call cap of 64, and a token-limit setting of 0. These values are present for every C/D/E and γ -sensitivity observation and are summarized in Table 8; equivalent configuration evidence is unavailable for B. A model name does not uniquely identify the inference version. Uncertainty in effective




Figure 5: A-E code-branch trajectories (median and interquartile range).

Productivity estimates and scheduling multipliers: three E cases
 Selection: lower-terminal coverage run, then two most observed states





Figure 6: E-arm pre-update probabilities (top) and scheduling multipliers (bottom).

Table 13: Branch endpoints (top) and AUCs (bottom), mean $\pm$ SD.

Target	A	B	C	D	E
LightFTP	71.0 $\pm$ 0.0 (n=13) 1758.5 $\pm$ 33.8	71.3 $\pm$ 0.9 (n=11) 1728.1 $\pm$ 36.5	80.6 $\pm$ 15.5 (n=10) 1987.7 $\pm$ 394.1	71.7 $\pm$ 2.3 (n=15) 1780.7 $\pm$ 45.2	73.9 $\pm$ 4.5 (n=10) 1825.0 $\pm$ 107.6
bftpd	486.0 $\pm$ 8.8 (n=10) 11868.7 $\pm$ 165.6	492.2 $\pm$ 14.1 (n=10) 12068.6 $\pm$ 323.8	498.9 $\pm$ 12.2 (n=10) 12082.0 $\pm$ 451.0	502.7 $\pm$ 7.5 (n=10) 12267.7 $\pm$ 387.9	503.7 $\pm$ 7.1 (n=10) 12416.6 $\pm$ 220.9
ProFTPD	5053.8 $\pm$ 97.4 (n=10) 122652.5 $\pm$ 1730.1	5205.9 $\pm$ 166.7 (n=10) 125217.7 $\pm$ 2989.1	5446.3 $\pm$ 194.7 (n=10) 130538.5 $\pm$ 4511.6	5428.1 $\pm$ 104.0 (n=12) 132539.2 $\pm$ 2868.9	5469.8 $\pm$ 215.7 (n=10) 130564.7 $\pm$ 4604.1
Pure-FTPD	1076.4 $\pm$ 89.6 (n=11) 23054.1 $\pm$ 2084.4	1109.3 $\pm$ 40.0 (n=10) 24976.6 $\pm$ 1126.8	1221.2 $\pm$ 48.9 (n=10) 28860.8 $\pm$ 1645.7	1223.3 $\pm$ 29.7 (n=10) 27863.9 $\pm$ 1496.9	1215.4 $\pm$ 63.7 (n=10) 27894.3 $\pm$ 1737.6
Exim	3798.8 $\pm$ 43.3 (n=10) 91083.0 $\pm$ 1031.5	3761.7 $\pm$ 60.4 (n=10) 91301.7 $\pm$ 1501.4	3896.1 $\pm$ 25.6 (n=10) 94880.1 $\pm$ 928.3	3914.9 $\pm$ 20.9 (n=10) 93659.0 $\pm$ 573.1	3858.3 $\pm$ 32.8 (n=10) 93123.6 $\pm$ 924.6
Live555	2893.6 $\pm$ 14.5 (n=10) 71127.4 $\pm$ 517.3	2882.4 $\pm$ 22.0 (n=10) 71224.8 $\pm$ 422.5	3075.7 $\pm$ 34.4 (n=10) 73511.0 $\pm$ 562.9	3067.1 $\pm$ 29.5 (n=12) 75436.1 $\pm$ 712.2	3054.2 $\pm$ 20.5 (n=10) 72378.5 $\pm$ 1049.8
Kamailio	10002.6 $\pm$ 144.2 (n=10) 242853.7 $\pm$ 2772.6	10644.3 $\pm$ 96.5 (n=10) 257606.8 $\pm$ 2082.0	10800.6 $\pm$ 129.0 (n=10) 260260.9 $\pm$ 2942.4	10631.2 $\pm$ 472.9 (n=12) 254332.1 $\pm$ 9492.1	10830.0 $\pm$ 91.3 (n=10) 259377.2 $\pm$ 2649.6
Forked-daapd	2330.8 $\pm$ 60.0 (n=10) 54453.2 $\pm$ 1494.1	2169.6 $\pm$ 113.0 (n=10) 49393.2 $\pm$ 3335.9	2585.4 $\pm$ 86.2 (n=10) 61776.8 $\pm$ 1863.7	2670.2 $\pm$ 87.3 (n=12) 63646.7 $\pm$ 1977.0	2609.0 $\pm$ 96.6 (n=11) 61843.9 $\pm$ 1570.4
Lighttpd1	1940.7 $\pm$ 21.2 (n=10) 46909.4 $\pm$ 235.7	1884.1 $\pm$ 25.3 (n=10) 45075.8 $\pm$ 729.3	2145.0 $\pm$ 69.8 (n=10) 53003.8 $\pm$ 1613.7	2126.7 $\pm$ 15.7 (n=11) 51223.2 $\pm$ 621.1	2068.1 $\pm$ 44.5 (n=10) 50385.1 $\pm$ 941.3

Table 14: C/D/E admission records and exploratory reward diagnostics.

Metric	C	D	E	Denominator / interpretation
Candidate records	1,326	2,227	746	Within-run distinct candidates.
Trial records	1,326	2,226	745	Within-run distinct trials.
Reject / durable labels	0 / 1,326	2,203 / 23	736 / 9	Reported trial outcomes.
First failed $P/U/R$	0 / 1,208 / 0	0 / 2,071 / 0	0 / 698 / 0	Trial denominator; first false P , U , or R .
Direct durable rate	100.0% (1,326/1,326)	1.0% (23/2,227)	1.2% (9/746)	Durable labels / generated candidates.
State-only provisional rate	disabled	0.0% (0/2,227)	0.0% (0/746)	All candidates; disabled for C.
Provisional conversion / expiration	disabled	none (0 admitted)	none (0 admitted)	Provisional admissions; retain pending/censored cases.
Trial latency (ms)	439.2 $\pm$ 405.9 ($n = 67$)	403.6 $\pm$ 574.6 ($n = 81$)	806.5 $\pm$ 554.0 ($n = 62$)	Within-run trial means; contributing runs only.
Recorded code-edge gain / episode	0.009 ($n = 3,986$)	0.003 ($n = 6,604$)	0.016 ($n = 1,474$)	265-mutation subset; recorded coverage gain.
BS / ECE / AP	0.240 / 0.428 / 0.012	0.192 / 0.373 / 0.005	0.315 / 0.516 / 0.038	265-mutation subset; pooled proxy reward; see Table 12.

Figure 7: C/D/E reported dispositions (left) and first $P/U/R$ failures (right).

conditions and resource measurement limits interpretation, so these values do not establish the shared confirmatory settings required of B–E.

Candidate-associated token counts may exclude other request types or repeated attempts; their sum therefore need not

equal total campaign cost. Low-activation targets, high-rejection targets, token outliers, saturation, and regressions must remain in the final analysis. Low activation may reflect productive mutation or a restrictive trigger; high rejection may reflect either useful gating or false negatives. Both need mechanism evidence.

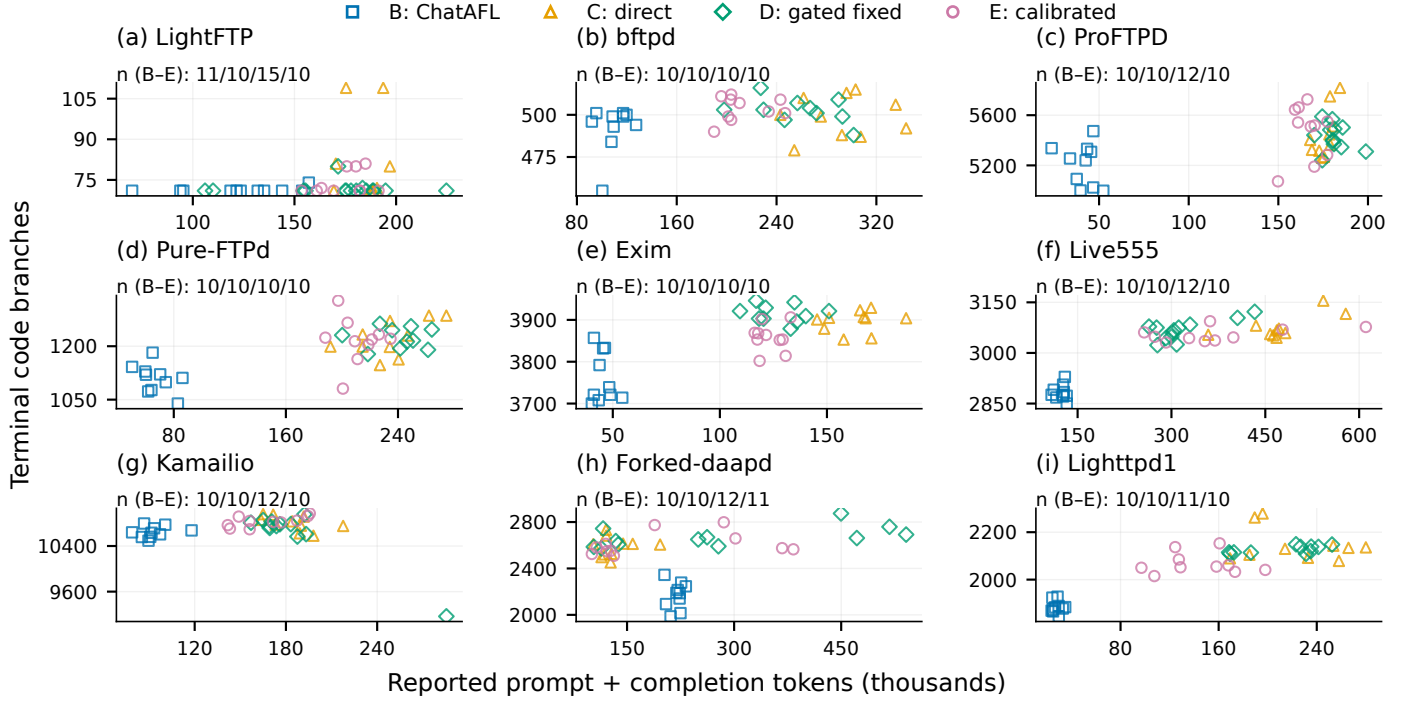


Figure 8: Code-branch endpoints versus reported tokens, by target and arm.

Table 15: D-arm cost ranges; total tokens = input + output ($k = 1,000$).

Target	Model calls	Total tokens (k)
LightFTP	59–161	106.0–224.7
bftpd	199–248	198.0–301.8
ProFTPD	188–192	170.1–198.8
Pure-FTPD	207–242	200.2–264.1
Exim	106–123	109.4–151.0
Live555	123–189	264.1–433.1
Kamilio	56–108	157.0–285.5
Forked-daapd	52–163	103.4–540.9
Lighttpd1	131–187	168.5–252.5

Figures 8 and 9 plot reported token usage and model calls against terminal code branches for 376 B–E runs, using target-specific scales. Tokens combine prompt and completion counts and are displayed in thousands (k). Comparable AFLNet usage measurements are unavailable; calls measure neither token expenditure nor CPU time. These views reveal dispatch and usage ranges without estimating a causal return per token or call. Measured CPU-hours and controlled A–E CVE recall remain unavailable. The case-level replays in Section 8 do not supply those cost–performance axes.

Repair remains an optional, unquantified procedure. No repair outcomes were recorded, and the study reports no repair ablation, effectiveness estimate, or repair-specific cost comparison.

8. Vulnerability Replay Evidence and Rediscovery Protocol

The vulnerability study comprises four cases: three associated with CVE identifiers and one SMARTPL availability

failure without a CVE. Each case has a report and a triggering input; three have direct execution evidence, and ProFTPD also has a reported patched-version comparison. Tables 16 and 17 summarize these observations. Case identities, versions, and replay counts follow the original reports; incomplete patch and experimental-configuration evidence prevents independent verification of every version pair.

8.1. Case identity and evidence provenance

We distinguish a direct execution observation (L) from a replay result stated only in a case report (R). A report can contain additional trials not represented by the available execution evidence. Duplicate diagnostic observations do not constitute independent trials. Input identity and byte length are retained in the supporting evidence. This retrospective analysis involves no additional target executions. Table 16 identifies the reported versions, triggering conditions, and comparative evidence.

For ProFTPD, the reported defect is an out-of-bounds read during FTP command processing, supported by an ASan heap-buffer-overflow observation. The 60,006-byte reproducer caused a one-byte read beyond a 65,568-byte region. The oracle uses a 65,535-byte command buffer; the report states that the default 512-byte configuration does not expose the same error to ASan. The result is therefore conditional on this setting. The full replay in Table 17 compares the campaign’s base version with the same version incorporating only the minimal upstream security correction. In this pair, 172 of 233 campaign crash inputs satisfy the FTP command-processing memory-error criterion in the vulnerable version, whereas none of 242 activate ASan in the patched version. This establishes differential replay for the verified pair; the three-trial patch comparison remains separate.

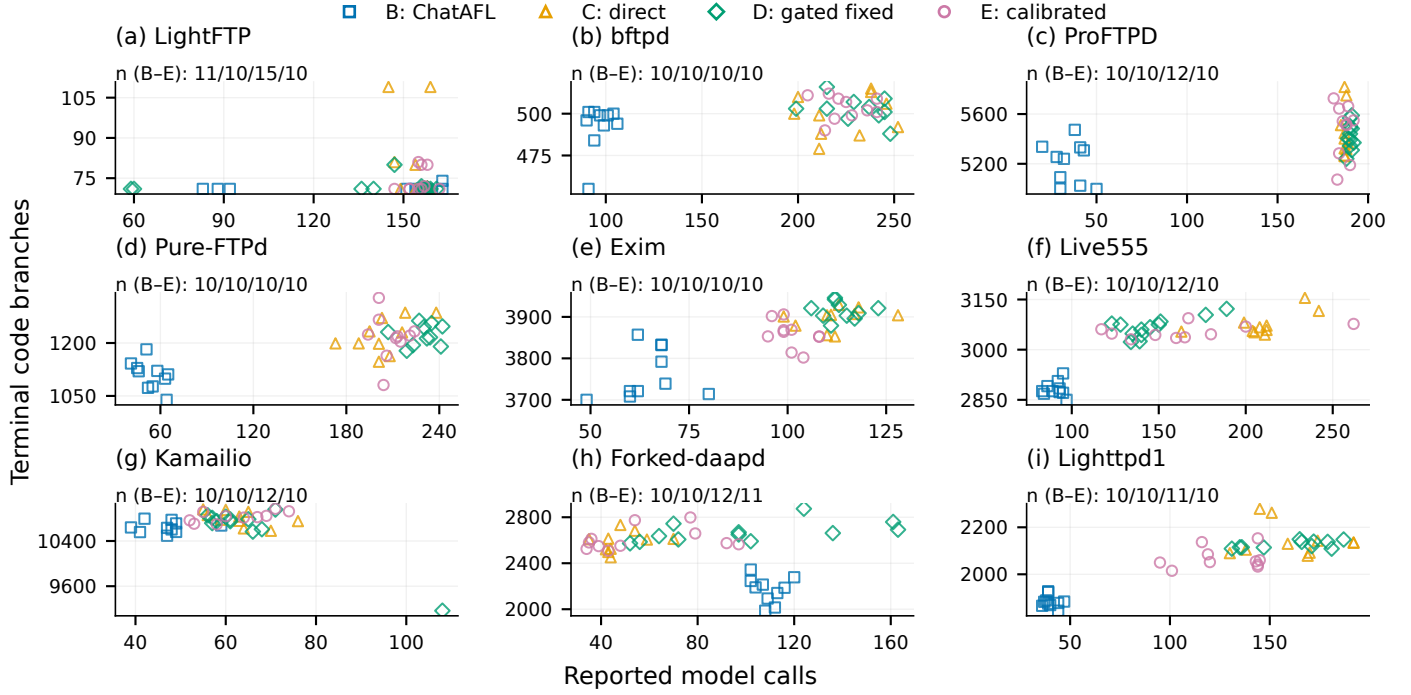


Figure 9: Code-branch endpoints versus reported calls, by target and arm.

Table 16: Vulnerability cases and supporting evidence (L: direct observation; R: report).

Case / interface	Reported vulnerable version	Trigger evidence / oracle	Patch or control evidence
CVE-2023-51713 ProFTPD / FTP	Reported vulnerable version; 65,535-byte buffer.	60,006-byte input; ASan heap-buffer-overflow (L).	Security patch: three rejections, no ASan (L).
CVE-2026-38998 LIVE555 / RTSP	Campaign: 2023.05.10; report: 2026.02.26 (R).	728-byte public PoC; use-after-free (L).	Local patch: 1/1 triggers (R); upstream comparison unavailable.
CVE-2026-39863 Kamailio / SIP over TCP	Reported vulnerable version; TCP required.	278-byte input; parsed Content-Length -10 (L).	Patched: 0/253 negative-value outcomes; security check observed (L).
SMARTPL availability case forked-daapd / HTTP	Reported version 27.2.	148-byte input; unresponsive service; process active (R).	No patched replay/CVE; removal reported in 28.4.

Reported input-specific replays yielded 1/1 seed, 1/1 fuzzer-crash, and 3/3 batch-crash reproductions; the separate patch comparison yielded 0/3 using the same input and buffer.

For LIVE555, direct diagnostic evidence supports a use-after-free triggered by the public PoC. A duplicated observation is treated as one diagnostic signature rather than independent replication. The report attributes fuzzer-generated inputs to the same memory-lifetime failure. It gives 1/1 public-PoC reproduction, 3/3 repeats for each of three fuzzer inputs, an 8/9 aggregate, and 1/1 persistence after a local patch. The fuzzer-generated reproducer, complete memory-lifetime evidence, and individual trial outcomes are unavailable. The public PoC supports replay validation, but does not demonstrate independent rediscovery by the evaluated controller or independently establish cross-version CVE attribution.

For Kamailio, a negative parsed Content-Length supports the reported integer-overflow condition. The trigger requires TCP and is not assessed by the UDP-only benchmark condition.

The complete-input comparison in Table 17 evaluates every retained input with an oversized Content-Length against two versions that differ only by the upstream security check. Among these inputs, 102 of 253 produce a negative-length outcome in the vulnerable version and none do in the patched version; the security check is observed in 193 of 253 patched-version trials. Reported input batches remain separate: 1/1 seed and 3/4, 19/20, 19/20, and 50/51 selected inputs. The comparison using a local security check is also separate. These inputs were generated during fuzzing, so the analysis measures predicate yield among retained inputs rather than independent rediscovery.

For SMARTPL, the 148-byte input reportedly causes lexical-analysis errors and service unresponsiveness while the process remains active. Reported repeats comprise 5/5 advisory trials, 3/3 trials with a 48-message seed, and a separate 1/1 replay. None of the five recorded campaign hangs reproduces this condition; a CVE and verified patched counterpart remain unavailable. Removal of the affected parsing component in version 28.4 re-

Table 17: Replay outcomes and controls; original denominators retained.

Case	Replay outcome	Control outcome	Evidence limit
CVE-2023-51713	172/233 ASan triggers (L).	Patched: 0/242 ASan activations (L).	Input replay; campaign recall unmeasured.
CVE-2026-38998	82/88 server terminations (L).	Local patch: 1/1 triggers (R).	No sanitizer or verified patched version.
CVE-2026-39863	102/253 negative-value outcomes (L).	Patched: 0/253; security check: 193/253 (L).	Selected queue entries; campaign recall unmeasured.
SMARTPL availability case	Advisory: 5/5; seed: 3/3; separate: 1/1 (R).	Campaign hangs: 0/5 matching outcomes (R).	Known-input replay; no patch or CVE.

mains unverified by patched-version replay.

8.2. *Replay outcomes and their denominators*

Table 17 summarizes the replay outcomes; input-specific reported batches remain separate in the case descriptions. A numerator counts replay attempts or selected queue inputs satisfying that case’s predicate; the denominator is the number tested. These are not independent 12-hour fuzzer runs, and the main experiment’s nominal $N = 10$ is not substituted for the actual replay denominator. Selected batches, per-input repeats, and aggregate counts may overlap; they are not pooled into one success rate. In particular, 19/20 selected inputs is not 95% campaign recall, and 8/9 LIVE555 replays is below the prospective 95% stability threshold.

Complete-input replays assess the retained inputs. For ProFTPD, every crash-associated input from the two 24-hour LoopFuzz campaigns (233 inputs during active fuzzing: 121 and 112 per run; 9 during termination) was replayed once against a verified version pair with a common base version and only the minimal upstream security correction. In the vulnerable version, 172 of the 233 inputs from active fuzzing satisfy the ASan heap-buffer-overflow criterion in FTP command processing. All 172 share one diagnostic signature; no alternative appears. The patched version shows no ASan activation for any of the 242 inputs. The 61 non-reproducing inputs and the 9 termination-phase inputs may depend on campaign states not reconstructed by isolated replay. Original campaign diagnostics are insufficient; outcome classification therefore relies on the controlled version-pair replay.

For LIVE555, all 88 campaign crash inputs were replayed against the tested version without sanitizer checks: 82 terminate the server and 6 do not. Without a patched comparison or memory-error diagnosis, these results establish reproduction in that version only, not attribution to the use-after-free.

For Kamailio, all 253 retained inputs with Content-Length $\geq 2^{31}$ were evaluated against the verified vulnerable/patched pair: 102 produce a negative-length outcome in the vulnerable version and none in the patched version. The two runs contributed 77 and 176 inputs, with yields of 36/77 and 66/176. Overflow values include -2147483648 , -1 , and -954437177 ; 151/253 inputs do not produce the outcome. The roles of interaction order and request interpretation remain unresolved. These comparisons characterize input replay, without arm-level recall, time-to-trigger, or cost denominators.

Additional reported crashes and campaigns cannot be linked to the primary benchmark and ablation observations. Crash totals do not count unique vulnerabilities, and durations such as 60 or 1,540 minutes do not specify time to first trigger. Associations among reproducing inputs, verified A–E conditions, model information exposure, and independent root causes are incomplete. The replay outcomes therefore do not establish arm-specific discovery probabilities, novel CVEs, or admission/calibration gains.

8.3. *Evidence limits for controlled rediscovery*

Confirmatory rediscovery requires a vulnerable version and the same version with a minimal security patch, holding dependencies, instrumentation, settings, and input compatibility constant. Each case needs a network-reachable trigger, independent reached/triggered criteria, reproducible initial conditions, and a validated patch comparison. The proposed six-case pilot across at least three protocols, followed by expansion toward twelve, remains incomplete. Eligibility requires at least 95% observed oracle repeatability in a declared validation set, no activation after patching, reproducible experimental preparation, and information separation. The four cases do not constitute six qualified version pairs; a small all-success replay batch does not establish an underlying success probability of at least 95%.

Information-controlled campaigns must exclude CVE identifiers, advisories, public PoCs, and patch contents from model inputs and campaign seeds, while retaining anonymized case-to-oracle correspondence separately. The public PoC and known triggers in this collection are validation inputs; their replay is not evidence of information-controlled rediscovery. Source-level trigger conditions, sanitizer reports with patch validation, and protocol invariants distinguish reaching relevant code from actually triggering a defect. Experiments require isolated target networks and reset between trials; these are protocol requirements, not properties inferred from the reports.

For each CVE and arm, the confirmatory design specifies ten independent runs and a 12-hour deadline before execution. Deadline recall, time-to-trigger, reached-but-not-triggered counts, patched-target activation, executions/coverage/tokens before trigger, and recall per CPU-hour or 100,000 tokens remain unavailable for the controlled A–E comparison. Nontriggers require right-censoring records, and interrupted runs require separate failure accounting. The available replay batches do not

define those risk sets, so Kaplan–Meier estimates or arm-level security-effectiveness claims are not reported.

9. Discussion

9.1. *What calibration can and cannot repair*

The controller calibrates the instrumental utility of an observed state proxy; it does not recover or verify semantic protocol state. If one response identifier aliases several server states, a single productivity estimate averages them. Reweighting the identifier cannot reconstruct lost distinctions. Richer response interpretation or internal-state observation addresses different aspects but requires target-specific knowledge and measurement access.

Low productivity can arise from a poor response proxy, saturated coverage, restrictive grammar, inadequate seeds, or insufficient energy. Independent reward prevents IPSM novelty from certifying itself but does not automatically diagnose these causes. The boundary analysis considers aliasing, limited initial evidence, campaign phase, and target-specific regressions.

A fixed provisional horizon can discard a bridge whose payoff requires more descendants. Shadow validation and horizon sensitivity would help quantify this risk, but neither is available here; expired or rejected candidates are not assumed worthless. Likewise, U/R may exclude useful error-handling inputs. Independent oracle retention is a design requirement, while its experimental adequacy and admission false negatives remain unverified.

9.2. *Cost, nonstationarity, and model variability*

Selection-indexed discounting can retain stale evidence for rarely selected states. Priors and the exploration floor preserve opportunities without guaranteeing fast identification. When code coverage saturates, nearly every reward may be zero, and calibration may offer little benefit.

Model-service nondeterminism persists despite identical prompts and nominal parameters. A shared model name may correspond to different inference versions, repeated requests affect costs, and target responses alter later prompts. Equal-policy/equal-cap comparisons control declared resources; actual usage and cost-matched analysis characterize the remaining tradeoffs. Replay assesses decisions under fixed conditions rather than reproducing an adaptive campaign.

9.3. *Evidence required for confirmatory claims*

A confirmatory calibration claim requires nondegenerate predictions, reduced investment in branch-unproductive Forked-daapd states when present, and no systematic suppression of Lighttpd1 solely because of limited IPSM novelty. It also requires acceptable throughput on at least three pilot targets and better predictive diagnostics than the shared prequential base-rate reference. D-condition productivity estimates provide an additional diagnostic; the uncalibrated queue-yield statistic is not a probability. An E–D difference must be supported by a plausible measured mechanism. These are prospective eligibility criteria rather than completed significance tests.

In this study, calibration is exploratory, repair is unquantified, and the empirical scope is an integrated-system boundary analysis. Missing validation is not a measured null effect. An admission or calibration claim requires independent supporting evidence; a rediscovery claim requires controlled campaigns rather than replay alone. Parameters and analysis decisions must be fixed before confirmatory observation.

The planned confirmatory evidence comprises D–C and E–D estimates, complete candidate-level measurements, at least nine targets, twenty runs per core arm, and a reproducible historical benchmark of approximately eight to twelve vulnerabilities. The four replay cases do not meet that benchmark requirement. These are intended sample sizes and coverage criteria, not achieved results.

10. Threats to Validity

Table 18 summarizes the controls and remaining evidence limits.

The analysis is retrospective. Equivalence between the assessed system and every experimental version is unverified, so methodological concerns and observed outcomes retain separate evidential status. A nonsignificant adjusted comparison does not establish equivalence, and favorable endpoints do not identify a mechanism when duration, resources, or effective procedures differ. Admission, calibration, and rediscovery each require independent supporting evidence.

Benchmark properties can change the relative ranking of fuzzers: initial seed coverage and execution speed are empirically consequential covariates [50]. Consequently, a favorable endpoint under unequal experimental conditions cannot isolate the admission or scheduling effect. Mutation-based assessment offers a complementary fault-oriented perspective [51]; here, code coverage and IPSM growth are kept separate from independently verified vulnerability outcomes.

11. Conclusion

We present a LoopFuzz controller design with retrospective boundary analysis. Admission separates code progress from IPSM novelty and bounds provisional validation for state-only proposals. Productivity calibration remains exploratory and does not recover semantic protocol states.

The available observations describe coverage, model usage, admission labels, and proxy-reward behavior under incompletely verified experimental conditions. Four vulnerability cases add replay evidence with differing patch support. These measurements do not establish reduced queue pollution, improved downstream productivity, greater cost efficiency, or improved rediscovery. The contribution is a controller specification and validation assessment; repair remains unquantified.

CRedit authorship contribution statement

Author contribution roles remain to be confirmed by all authors before submission.

Table 18: Validity threats, required controls, and evidence limits.

Dimension	Required control	Current limitation
Internal	Admission-only C–D; scheduling-only E–D; repair off; matched initial states.	Missing observations; uncertain decision order.
Construct	Separate code branches, IPSM state edges, bitmap counts, and oracle outcomes.	Proxies incompletely measure intended criteria.
Statistical	Repeated runs; test families; effect sizes; clustered/hierarchical CIs; censoring.	Unequal durations; no confirmatory effects.
Selection	Intention-to-run accounting; prespecified repeats after experimental environment failure.	Unobserved planned runs; independence uncertain.
Reproducibility	Match versions, settings, seeds, prompts, and measurements.	Unresolved model versions; incomplete provenance.
External	Multiple targets/protocols; boundary cases; observability assumptions.	Encrypted/opaque binary protocols untested.
Security	Minimal patch pairs; independent triggers; information isolation.	Controlled A–E/patch-pair evidence incomplete; pretraining leakage unresolved.

Declaration of competing interest

The competing-interest declaration remains to be confirmed by all authors.

Funding

Funding information remains to be confirmed by all authors.

Data availability

The 642 experimental records and four vulnerability case studies are accompanied by trial-level measurements, version and input provenance, and numerical figure data. Public access and persistent identifiers remain unspecified; exact reproduction of model-service responses is not claimed.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

OpenAI Codex assisted language revision, organization, figure preparation, and consistency assessment. The authors remain responsible for scientific accuracy and the final disclosure.

Appendix A. Complete descriptive results

Table A.1 reports all observations underlying the descriptive comparisons, using the same values and sample SDs as the main results. Sample sizes are shown once per row; unavailable target–arm conditions are not imputed. Conditions are indexed by arm and target. Arms A–E follow the main text; E_{99} : $\gamma = 0.99$; E_{100} : $\gamma = 1.0$.

1: LightFTP 4: Pure-FTPD 7: Kmailio
2: bftpd 5: Exim 8: Forked-daapd
3: ProFTPD 6: Live555 9: Lighttpd1

AUC: branch-hours; IPSM: state edges. Reported status: C, completed; X, interruption near limit; S, abnormal termination; O, memory exhaustion. Causes are unverified. Missing combinations are unobserved, not zero.

ID	<i>n</i>	Branches	AUC	IPSM Status
A1	13	71.0±0.0	1758.5±33.8	187.3±11.3 C:13
A2	10	486.0±8.8	11868.7±165.6	163.1±15.4 C:10
A3	10	5053.8±97.4	122652.5±1730.1	190.5±27.9 C:10
A4	11	1076.4±89.6	23054.1±2084.4	236.5±26.7 C:11
A5	10	3798.8±43.3	91083.0±1031.5	88.8±5.2 C:10
A6	10	2893.6±14.5	71127.4±517.3	96.1±8.6 C:10
A7	10	10002.6±144.2	242853.7±2772.6	109.4±9.7 C:10
A8	10	2330.8±60.0	54453.2±1494.1	21.6±2.1 C:10
A9	10	1940.7±21.2	46909.4±235.7	17.7±2.5 C:10
B1	11	71.3±0.9	1728.1±36.5	195.1±14.6 C:11
B2	10	492.2±14.1	12068.6±323.8	167.1±12.5 C:10
B3	10	5205.9±166.7	125217.7±2989.1	150.0±18.5 C:10
B4	10	1109.3±40.0	24976.6±1126.8	186.0±31.5 C:10
B5	10	3761.7±60.4	91301.7±1501.4	115.8±13.5 C:10
B6	10	2882.4±22.0	71224.8±422.5	143.8±8.8 C:10
B7	10	10644.3±96.5	257606.8±2082.0	106.5±6.8 C:10
B8	10	2169.6±113.0	49393.2±3335.9	17.7±4.9 C:10
B9	10	1884.1±25.3	45075.8±729.3	14.3±3.1 C:10
C1	10	80.6±15.5	1987.7±394.1	210.6±8.3 C:10
C2	10	498.9±12.2	12082.0±451.0	217.9±13.8 C:6,X:4
C3	10	5446.3±194.7	130538.5±4511.6	300.7±29.1 X:9,C:1
C4	10	1221.2±48.9	28860.8±1645.7	328.5±22.5 C:7,X:3
C5	10	3896.1±25.6	94880.1±928.3	150.9±9.5 C:10
C6	10	3075.7±34.4	73511.0±562.9	166.8±9.3 C:10
C7	10	10800.6±129.0	260260.9±2942.4	139.2±9.8 C:9,X:1
C8	10	2585.4±86.2	61776.8±1863.7	22.3±2.2 C:8,S:2
C9	10	2145.0±69.8	53003.8±1613.7	32.1±5.0 C:10
D1	15	71.7±2.3	1780.7±45.2	207.0±17.0 C:15
D2	10	502.7±7.5	12267.7±387.9	226.7±14.0 C:10
D3	12	5428.1±104.0	132539.2±2868.9	305.3±20.3 C:12
D4	10	1223.3±29.7	27863.9±1496.9	321.4±11.1 C:10
D5	10	3914.9±20.9	93659.0±573.1	152.7±9.3 C:10
D6	12	3067.1±29.5	75436.1±712.2	170.1±11.1 C:12
D7	12	10631.2±472.9	254332.1±9492.1	128.2±26.6 C:12
D8	12	2670.2±87.3	63646.7±1977.0	22.8±2.9 C:12
D9	11	2126.7±15.7	51223.2±621.1	28.9±2.7 C:11
E1	10	73.9±4.5	1825.0±107.6	210.9±9.3 C:10
E2	10	503.7±7.1	12416.6±220.9	215.8±7.5 C:9,O:1
E3	10	5469.8±215.7	130564.7±4604.1	307.8±21.3 X:5,C:5

Table A.1: Target–arm coverage (mean ± SD; actual *n*, nominal *N* = 10).

Table A.1 (continued)

ID	<i>n</i>	Branches	AUC	IPSM Status
E4	10	1215.4±63.7	27894.3±1737.6	311.6±22.4 X:5,C:5
E5	10	3858.3±32.8	93123.6±924.6	138.6±14.3 C:10
E6	10	3054.2±20.5	72378.5±1049.8	162.1±4.4 C:10
E7	10	10830.0±91.3	259377.2±2649.6	130.9±9.0 C:10
E8	11	2609.0±96.6	61843.9±1570.4	21.8±1.7 C:11
E9	10	2068.1±44.5	50385.1±941.3	27.1±4.4 C:10
E ₉₉ 1	10	71.9±2.8	1785.4±78.3	210.6±12.4 C:10
E ₉₉ 2	10	496.8±9.2	12210.9±353.6	224.0±9.8 C:8,X:2
E ₉₉ 3	10	5504.6±132.2	134386.6±3705.3	310.5±19.2 C:9,X:1
E ₉₉ 4	10	1219.4±26.3	28302.9±964.5	310.5±19.8 C:10
E ₉₉ 5	10	3856.3±26.9	94717.9±787.4	158.9±13.5 X:3,C:7
E ₉₉ 6	7	3091.9±19.9	75920.1±441.3	161.3±6.8 C:7
E ₉₉ 7	9	10852.9±93.4	265216.0±2511.7	137.7±6.5 C:8,S:1
E ₉₉ 8	9	2517.6±26.2	60824.3±1014.6	21.2±2.4 C:8,S:1
E ₉₉ 9	6	2073.3±28.7	50599.8±669.8	28.5±7.1 C:6
E ₁₀₀ 1	10	72.8±4.7	1811.5±143.3	213.1±9.4 C:10
E ₁₀₀ 2	10	500.3±5.7	12417.2±192.6	223.1±10.6 C:9,X:1
E ₁₀₀ 3	10	5429.7±158.7	131751.6±3888.9	304.4±22.5 X:2,C:8
E ₁₀₀ 4	10	1181.2±23.6	26741.3±1359.4	302.8±16.7 X:1,C:9
E ₁₀₀ 5	10	3837.1±33.1	93785.6±875.9	155.9±12.6 C:10
E ₁₀₀ 6	10	3088.5±16.5	75880.3±914.8	162.4±6.3 C:9,S:1
E ₁₀₀ 7	11	10851.4±122.1	264236.5±3000.1	130.5±10.5 C:11
E ₁₀₀ 8	10	2525.4±45.5	61039.2±1214.6	21.6±2.7 C:10
E ₁₀₀ 9	10	2056.0±62.1	50356.8±1354.3	27.0±2.8 C:10

References

- [1] V.-T. Pham, M. Böhme, A. Roychoudhury, AFLNet: A greybox fuzzer for network protocols, in: 2020 IEEE 13th International Conference on Software Testing, Verification and Validation (ICST), 2020, pp. 460–465. doi:10.1109/ICST46399.2020.00062.
- [2] R. Natella, StateAFL: Greybox fuzzing for stateful network servers, Empirical Software Engineering 27 (2022). doi:10.1007/s10664-022-10233-3, art. no. 191.
- [3] R. Meng, M. Mirchev, M. Böhme, A. Roychoudhury, Large language model guided protocol fuzzing, in: 31st Annual Network and Distributed System Security Symposium (NDSS 2024), 2024, pp. 1–17. doi:10.14722/ndss.2024.24556.
- [4] S. Qin, F. Hu, Z. Ma, B. Zhao, T. Yin, C. Zhang, NSFuzz: Towards efficient and state-aware network service fuzzing, ACM Transactions on Software Engineering and Methodology 32 (2023) 1–26. doi:10.1145/3580598.
- [5] J. Ba, M. Böhme, Z. Mirzamomen, A. Roychoudhury, Stateful greybox fuzzing, in: 31st USENIX Security Symposium (USENIX Security 2022), 2022, pp. 3255–3272. URL: <https://www.usenix.org/conference/usenixsecurity22/presentation/ba>.
- [6] R. Meng, V.-T. Pham, M. Böhme, A. Roychoudhury, AFLNet Five Years Later: On Coverage-Guided Protocol Fuzzing, IEEE Transactions on Software Engineering 51 (2025) 960–974. doi:10.1109/tse.2025.3535925.
- [7] C. McMahon Stone, S. L. Thomas, M. Vanhoef, J. Henderson, N. Bailluet, T. Chothia, The Closer You Look, The More You Learn: A Grey-box Approach to Protocol State Machine Learning, in: Proceedings of the 2022 ACM SIGSAC Conference on Computer and Communications Security, 2022, pp. 2265–2278. doi:10.1145/3548606.3559365.
- [8] S. Schumilo, C. Aschermann, A. Jemmett, A. Abbasi, T. Holz, Nyx-net: network fuzzing with incremental snapshots, in: Proceedings of the Seventeenth European Conference on Computer Systems, 2022, pp. 166–180. doi:10.1145/3492321.3519591.
- [9] A. Andronidis, C. Cadar, SnapFuzz: high-throughput fuzzing of network applications, in: Proceedings of the 31st ACM SIGSOFT International Symposium on Software Testing and Analysis, 2022, pp. 340–351. doi:10.1145/3533767.3534376.
- [10] N. Bars, M. Schloegel, N. Schiller, L. Bernhard, T. Holz, No Peer, no Cry: Network Application Fuzzing via Fault Injection, in: Proceedings of the 2024 on ACM SIGSAC Conference on Computer and Communications Security, 2024, pp. 750–764. doi:10.1145/3658644.3690274.
- [11] W.-C. Kao, Y.-C. Chen, Y.-S. Lin, Y.-C. Yang, C.-Y. Li, C.-Y. Huang, BLuEMan: A stateful simulation-based fuzzing framework for Open-Source RTOS bluetooth low energy protocol stacks, in: 34th USENIX Security Symposium (USENIX Security 25), USENIX Association, Seattle, WA, 2025, pp. 6259–6278. URL: <https://www.usenix.org/conference/usenixsecurity25/presentation/kao>.
- [12] Y. Zou, J.-J. Bai, Z.-M. Jiang, M. Zhao, D. Zhou, Blackbox Fuzzing of Distributed Systems with Multi-Dimensional Inputs and Symmetry-Based Feedback Pruning, in: Proceedings 2025 Network and Distributed System Security Symposium, 2025. doi:10.14722/ndss.2025.241912.
- [13] P. Srivastava, M. Payer, Gramatron: effective grammar-aware fuzzing, in: Proceedings of the 30th ACM SIGSOFT International Symposium on Software Testing and Analysis, 2021, pp. 244–256. doi:10.1145/3460319.3464814.
- [14] D. Wang, Y. Li, Z. Zhang, K. Chen, CarpetFuzz: Automatic program option constraint extraction from documentation for fuzzing, in: 32nd USENIX Security Symposium (USENIX Security 2023), 2023, pp. 1919–1936. URL: <https://www.usenix.org/conference/usenixsecurity23/presentation/wang-dawei>.
- [15] R. Dutra, R. Gopinath, A. Zeller, FormatFuzzer: Effective Fuzzing of Binary File Formats, ACM Transactions on Software Engineering and Methodology 33 (2024) 1–29. doi:10.1145/3628157.
- [16] N. Bars, M. Schloegel, T. Scharnowski, N. Schiller, T. Holz, Fuzztruction: Using fault injection-based fuzzing to leverage implicit domain knowledge, in: 32nd USENIX Security Symposium (USENIX Security 23), USENIX Association, Anaheim, CA, 2023, pp. 1847–1864. URL: <https://www.usenix.org/conference/usenixsecurity23/presentation/bars>.
- [17] C. Lyu, J. Xu, S. Ji, X. Zhang, Q. Wang, B. Zhao, G. Pan, W. Cao, P. Chen, R. Beyah, MINER: A hybrid Data-Driven approach for REST API fuzzing, in: 32nd USENIX Security Symposium (USENIX Security 23), USENIX Association, Anaheim, CA, 2023, pp. 4517–4534. URL: <https://www.usenix.org/conference/usenixsecurity23/presentation/lyu>.
- [18] J. Wang, L. Yu, X. Luo, LLMIF: Augmented large language model for fuzzing IoT devices, in: 2024 IEEE Symposium on Security and Privacy (SP), 2024, pp. 881–896. doi:10.1109/sp54263.2024.00211.
- [19] C. S. Xia, M. Paltenghi, J. L. Tian, M. Pradel, L. Zhang, Fuzz4All: Universal fuzzing with large language models, in: Proceedings of the IEEE/ACM 46th International Conference on Software Engineering, 2024, pp. 1–13. doi:10.1145/3597503.3639121.
- [20] Y. Deng, C. S. Xia, C. Yang, S. D. Zhang, S. Yang, L. Zhang, Large Language Models are Edge-Case Generators: Crafting Unusual Programs for Fuzzing Deep Learning Libraries, in: Proceedings of the IEEE/ACM 46th International Conference on Software Engineering, 2024, pp. 1–13. doi:10.1145/3597503.3623343.
- [21] C. Yang, Y. Deng, R. Lu, J. Yao, J. Liu, R. Jabbarvand, L. Zhang, WhiteFox: White-Box Compiler Fuzzing Empowered by Large Language Models, Proceedings of the ACM on Programming Languages 8 (2024) 709–735. doi:10.1145/3689736.
- [22] D. Wang, G. Zhou, L. Chen, D. Li, Y. Miao, ProphetFuzz: Fully Automated Prediction and Fuzzing of High-Risk Option Combinations with Only Documentation via Large Language Model, in: Proceedings of the 2024 on ACM SIGSAC Conference on Computer and Communications Security, 2024,

- pp. 735–749. doi:10.1145/3658644.3690231.
- [23] R. Meng, G. J. Duck, A. Roychoudhury, Large Language Model Assisted Hybrid Fuzzing, *IEEE Transactions on Software Engineering* 52 (2026) 2367–2382. doi:10.1109/tse.2026.3694408.
 - [24] H. Xu, Y. Zhao, J. Yu, H. Wang, HGFuzzer: Directed Greybox Fuzzing via Large Language Model, *ACM Transactions on Software Engineering and Methodology* (2026). doi:10.1145/3841476, advance online publication, August 22, 2026.
 - [25] A. Borcharding, M. Giraud, I. Fitzgerald, J. Beyerer, The bandit’s states: Modeling state selection for stateful network fuzzing as multi-armed bandit problem, in: 2023 IEEE European Symposium on Security and Privacy Workshops (EuroS&PW), 2023, pp. 345–350. doi:10.1109/EuroSPW59978.2023.00043.
 - [26] S. Luo, A. Herrera, P. Quirk, M. Chase, D. C. Ranasinghe, S. S. Kanhere, Make out like a (multi-armed) bandit: Improving the odds of fuzzer seed scheduling with T-Scheduler, in: Proceedings of the 19th ACM Asia Conference on Computer and Communications Security, 2024, pp. 1463–1479. doi:10.1145/3634737.3637639.
 - [27] M. Böhme, V. J. M. Manès, S. K. Cha, Boosting fuzzer efficiency: An information theoretic perspective, in: Proceedings of the 28th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering, 2020, pp. 678–689. doi:10.1145/3368089.3409748.
 - [28] R. Ollando, S. Y. Shin, L. C. Briand, Learning-guided fuzzing for testing stateful SDN controllers, *ACM Transactions on Software Engineering and Methodology* 35 (2026) 1–45. doi:10.1145/3733717.
 - [29] D. She, A. Storek, Y. Xie, S. Kweon, P. Srivastava, S. Jana, FOX: Coverage-guided Fuzzing as Online Stochastic Control, in: Proceedings of the 2024 on ACM SIGSAC Conference on Computer and Communications Security, 2024, pp. 765–779. doi:10.1145/3658644.3670362.
 - [30] H. Zheng, J. Zhang, Y. Huang, Z. Ren, H. Wang, C. Cao, Y. Zhang, F. Toffalini, M. Payer, FISHFUZZ: Catch deeper bugs by throwing larger nets, in: 32nd USENIX Security Symposium (USENIX Security 23), USENIX Association, Anaheim, CA, 2023, pp. 1343–1360. URL: <https://www.usenix.org/conference/usenixsecurity23/presentation/zheng>.
 - [31] X. Zhu, E. Dai, X. Feng, S. Wang, X. Xia, S. Wen, K.-Y. Lam, Y. Xiang, WingMuzz: Blackbox Testing of IoT Protocols via Two-dimensional Fuzzing Schedule, in: 2025 40th IEEE/ACM International Conference on Automated Software Engineering (ASE), 2025, pp. 2580–2592. doi:10.1109/ase63991.2025.00212.
 - [32] X. Song, L. Pei, J. Wu, Y. Zeng, G. He, C. Zuo, X. Liu, Q. Zhao, S. Guo, ProtocolGuard: Detecting Protocol Non-compliance Bugs via LLM-guided Static Analysis and Dynamic Verification, in: Proceedings 2026 Network and Distributed System Security Symposium, 2026. doi:10.14722/ndss.2026.240521.
 - [33] T. Yang, Y. Qin, L. Zhang, Z. Fu, J. Chen, J. Wang, S. Zhao, Q. Li, R. Li, H. Wang, Y. Zhang, BSFuzzer: Context-Aware Semantic Fuzzing for BLE Logic Flaw Detection, in: Proceedings 2026 Network and Distributed System Security Symposium, 2026. doi:10.14722/ndss.2026.240094.
 - [34] K. Wang, J. Chen, P. Chen, J. Zhuge, J. Bai, H. Duan, Identifying Logical Vulnerabilities in QUIC Implementations, in: Proceedings 2026 Network and Distributed System Security Symposium, 2026. doi:10.14722/ndss.2026.231777.
 - [35] A. Hazimeh, A. Herrera, M. Payer, Magma: A ground-truth fuzzing benchmark, *Proceedings of the ACM on Measurement and Analysis of Computing Systems* 4 (2020) 1–29. doi:10.1145/3428334.
 - [36] R. Natella, V.-T. Pham, ProFuzzBench: A benchmark for stateful protocol fuzzing, in: Proceedings of the 30th ACM SIGSOFT International Symposium on Software Testing and Analysis, 2021, pp. 662–665. doi:10.1145/3460319.3469077.
 - [37] J. Metzman, L. Szekeres, L. Simon, R. Sprabery, A. Arya, FuzzBench: an open fuzzer benchmarking platform and service, in: Proceedings of the 29th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering, 2021, pp. 1393–1403. doi:10.1145/3468264.3473932.
 - [38] M. Böhme, L. Szekeres, J. Metzman, On the reliability of coverage-based fuzzer benchmarking, in: Proceedings of the 44th International Conference on Software Engineering, 2022, pp. 1621–1633. doi:10.1145/3510003.3510230.
 - [39] M. Schloegel, N. Bars, N. Schiller, L. Bernhard, T. Scharnowski, A. Crump, A. Ale-Ebrahim, N. Bissantz, M. Muench, T. Holz, SoK: Prudent evaluation practices for fuzzing, in: 2024 IEEE Symposium on Security and Privacy (SP), 2024, pp. 1974–1993. doi:10.1109/sp54263.2024.00137.
 - [40] W. R. Thompson, On the likelihood that one unknown probability exceeds another in view of the evidence of two samples, *Biometrika* 25 (1933) 285–294. doi:10.1093/biomet/25.3-4.285.
 - [41] J. Ounjai, V. Wüstholtz, M. Christakis, Green Fuzzer Benchmarking, in: Proceedings of the 32nd ACM SIGSOFT International Symposium on Software Testing and Analysis, 2023, pp. 1396–1406. doi:10.1145/3597926.3598144.
 - [42] C. Tantithamthavorn, S. McIntosh, A. E. Hassan, K. Matsumoto, An empirical comparison of model validation techniques for defect prediction models, *IEEE Transactions on Software Engineering* 43 (2017) 1–18. doi:10.1109/TSE.2016.2584050.
 - [43] C. Guo, G. Pleiss, Y. Sun, K. Q. Weinberger, On calibration of modern neural networks, in: Proceedings of the 34th International Conference on Machine Learning, volume 70 of *Proceedings of Machine Learning Research*, PMLR, 2017, pp. 1321–1330. URL: <https://proceedings.mlr.press/v70/guo17a.html>.
 - [44] J. Davis, M. Goadrich, The relationship between Precision-Recall and ROC curves, in: Proceedings of the 23rd International Conference on Machine Learning, 2006, pp. 233–240. doi:10.1145/1143844.1143874.
 - [45] B. Efron, Bootstrap methods: Another look at the jackknife, *The Annals of Statistics* 7 (1979) 1–26. doi:10.1214/aos/1176344552.
 - [46] H. B. Mann, D. R. Whitney, On a test of whether one of two random variables is stochastically larger than the other, *The Annals of Mathematical Statistics* 18 (1947) 50–60. doi:10.1214/aoms/1177730491.
 - [47] N. Cliff, Dominance statistics: Ordinal analyses to answer ordinal questions, *Psychological Bulletin* 114 (1993) 494–509. doi:10.1037/0033-2909.114.3.494.
 - [48] Y. Benjamini, Y. Hochberg, Controlling the false discovery rate: A practical and powerful approach to multiple testing, *Journal of the Royal Statistical Society Series B: Statistical Methodology* 57 (1995) 289–300. doi:10.1111/j.2517-6161.1995.tb02031.x.
 - [49] E. L. Kaplan, P. Meier, Nonparametric estimation from incomplete observations, *Journal of the American Statistical Association* 53 (1958) 457–481. doi:10.1080/01621459.1958.10501452.
 - [50] D. Wolff, M. Böhme, A. Roychoudhury, Fuzzing: On Benchmarking Outcome as a Function of Benchmark Properties, *ACM Transactions on Software Engineering and Methodology* 35 (2026) 1–26. doi:10.1145/3732936.
 - [51] P. Götz, B. Mathis, K. Hassler, E. Güler, T. Holz, A. Zeller, R. Gopinath, Systematic assessment of fuzzers using mutation analysis, in: 32nd USENIX Security Symposium (USENIX Security 23), USENIX Association, Anaheim, CA, 2023, pp. 4535–4552. URL: <https://www.usenix.org/conference/usenixsecurity23/presentation/gorz>.